

main

November 4, 2021

```
[1]: import torch
import torchvision
from torchvision import transforms, datasets
import matplotlib.pyplot as plt
import numpy as np
from torch.utils.data import Subset
from torch.utils.data import DataLoader
import random
import torch.nn as nn
import torch.nn.functional as F
from torch.optim import Adam
from sklearn.metrics import classification_report
from matplotlib import pyplot as plt
import time
from sklearn.metrics import confusion_matrix
from torch.utils.tensorboard import SummaryWriter
from torch.autograd import Variable
import cv2
%reload_ext tensorboard
```

c:\users\hp631\appdata\local\programs\python\python39\lib\site-packages\setuptools\distutils_patch.py:25: UserWarning: Distutils was imported before Setuptools. This usage is discouraged and may exhibit undesirable behaviors or errors. Please use Setuptools' objects directly or at least import Setuptools first.

```
warnings.warn(
```

```
[2]: def dataset_builder(size_img):

    #reading the train and test datasets and converting them to tensor
    #transforms.ToTensor will scale the images to 0-1 range
    #resize images using transforms.resize
    data_transform = transforms.Compose([
        transforms.Resize((size_img,size_img)),
        transforms.ToTensor()])
    train_dataset = datasets.ImageFolder(root='../data/train/img',
                                         transform=data_transform)
    train_dataloader = DataLoader(train_dataset,batch_size=len(train_dataset),
```

```

        shuffle=True,
        num_workers=4)

test_val_dataset=datasets.ImageFolder(root='../data/test/img',
                                       transform=data_transform)

#splitting test dataset into test+val dataset such that each class in test
→and validation dataset has
# 500 and 300 images respectively. Also, Shuffling is done before the sets
→are created.
dataset_size = len(test_val_dataset)
val_idx=[]
test_idx=[]
temp=[]
for i in range(0, 8000, 800):
    temp=random.sample(range(i,i+800),800)
    val_idx+=temp[0:300]
    test_idx+=temp[300:800]
val_dataset=Subset(test_val_dataset, val_idx)
test_dataset=Subset(test_val_dataset,test_idx)

#normalizing the datasets. first calculate mean and std for each channel in
→each training image and take
#the average of those values. Then normalize train,test,val datasets using
→these found values for training dataset.
inputs, classes = next(iter(train_dataloader))
#print("inputs",len(inputs))
mean_r=0
std_r=0
mean_g=0
std_g=0
mean_b=0
std_b=0
for img in inputs:
    mean_r+=(img[0].mean()).item()
    std_r+=(img[0].std()).item()
    mean_g+=(img[1].mean()).item()
    std_g+=(img[1].std()).item()
    mean_b+=(img[2].mean()).item()
    std_b+=(img[2].std()).item()
mean_R=mean_r/(len(inputs))
mean_G=mean_g/(len(inputs))
mean_B=mean_b/(len(inputs))
std_R=std_r/(len(inputs))
std_G=std_g/(len(inputs))
std_B=std_b/(len(inputs))

```

```

    #now normalize the datasets using these values of mean and std for each
    ↪channel
    data_transform_normalize = transforms.Compose([
        transforms.Resize((size_img,size_img)), transforms.ToTensor(),
        transforms.Normalize((mean_R,mean_G,mean_B),(std_R,std_G,std_B))])
    train_dataset_normalized = datasets.ImageFolder(root='../data/train/img',
                                                    transform=data_transform_normalize)

    val_dataset.dataset.transform = transforms.Compose([transforms.
    ↪Resize((size_img,size_img)), transforms.ToTensor(),
        transforms.Normalize((mean_R,mean_G,mean_B),(std_R,std_G,std_B))])

    test_dataset.dataset.transform = transforms.Compose([transforms.
    ↪Resize((size_img,size_img)), transforms.ToTensor(),
        transforms.Normalize((mean_R,mean_G,mean_B),(std_R,std_G,std_B))])

    return train_dataset_normalized, val_dataset, test_dataset

```

```

[3]: class LeNet5(nn.Module):
    def __init__(self):
        super().__init__()

        #initialize first layer:conv->relu-->maxpooling
        self.conv1=nn.Conv2d(in_channels=3, out_channels=6, kernel_size=5,
    ↪stride=1, padding=0)
        self.relu1=nn.ReLU()
        self.pool1=nn.MaxPool2d(kernel_size=2, stride=2, padding=0)

        #initialize second layer:conv->relu-->maxpooling
        self.conv2=nn.Conv2d(in_channels=6, out_channels=16, kernel_size=5,
    ↪stride=1, padding=0)
        self.relu2=nn.ReLU()
        self.pool2=nn.MaxPool2d(kernel_size=2, stride=2, padding=0)

        #initialize third layer:FC layer-->relu
        self.fc1=nn.Linear(in_features=16*5*5, out_features=120)
        self.relu3=nn.ReLU()

        #initialize fourth layer:FC layer-->relu
        self.fc2=nn.Linear(in_features=120, out_features=84)
        self.relu4=nn.ReLU()

        #initialize fifth layer:FC layer-->logSoftMax
        self.fc3=nn.Linear(in_features=84, out_features=10)
        self.logSoftmax = nn.LogSoftmax(dim=1)

```

```

def forward(self,x):
    #passing the input through first set of layers
    x=self.conv1(x)
    x=self.relu1(x)
    x=self.pool1(x)

    #passing output from first layer through second set of layers
    x=self.conv2(x)
    x=self.relu2(x)
    x=self.pool2(x)

    #passing output from second layer through third set of layers
    x=torch.flatten(x,1)
    x=self.fc1(x)
    x=self.relu3(x)

    #passing output from third layer through fourth set of layers
    x=self.fc2(x)
    x=self.relu4(x)

    #passing output from fourth layer through fifth set of layers
    x=self.fc3(x)
    output = self.logSoftmax(x)
    return output

```

```

[36]: class LeNet5_BN(nn.Module):
    def __init__(self):
        super().__init__()

        #initialize first layer:conv->BatchNormalization-->relu-->maxpooling
        self.conv1=nn.Conv2d(in_channels=3, out_channels=6, kernel_size=5,
↪stride=1, padding=0)
        self.bn1=nn.BatchNorm2d(6)
        self.relu1=nn.ReLU()
        self.pool1=nn.MaxPool2d(kernel_size=2, stride=2, padding=0)

        #initialize second layer:conv->BatchNormalization-->relu-->maxpooling
        self.conv2=nn.Conv2d(in_channels=6, out_channels=16, kernel_size=5,
↪stride=1, padding=0)
        self.bn2=nn.BatchNorm2d(16)
        self.relu2=nn.ReLU()
        self.pool2=nn.MaxPool2d(kernel_size=2, stride=2, padding=0)

        #initialize third layer:FC layer-->BatchNormalization-->relu
        self.fc1=nn.Linear(in_features=16*5*5, out_features=120)
        self.bn3=nn.BatchNorm1d(120)

```

```

self.relu3=nn.ReLU()

#initialize fourth layer:FC layer-->BatchNormalization-->relu
self.fc2=nn.Linear(in_features=120, out_features=84)
self.bn4=nn.BatchNorm1d(84)
self.relu4=nn.ReLU()

#initialize fifth layer:FC layer-->logSoftMax
self.fc3=nn.Linear(in_features=84, out_features=10)
self.logSoftmax = nn.LogSoftmax(dim=1)

def forward(self,x):
    #passing the input through first set of layers
    x=self.conv1(x)
    x=self.bn1(x)
    x=self.relu1(x)
    x=self.pool1(x)

    #passing output from first layer through second set of layers
    x=self.conv2(x)
    x=self.bn2(x)
    x=self.relu2(x)
    x=self.pool2(x)

    #passing output from second layer through third set of layers
    x=torch.flatten(x,1)
    x=self.fc1(x)
    x=self.bn3(x)
    x=self.relu3(x)

    #passing output from third layer through fourth set of layers
    x=self.fc2(x)
    x=self.bn4(x)
    x=self.relu4(x)

    #passing output from fourth layer through fifth set of layers
    x=self.fc3(x)
    output = self.logSoftmax(x)
    return output

```

```

[4]: train_dataset, val_dataset, test_dataset=dataset_builder(32)
print("train size",len(train_dataset))
print("test size",len(test_dataset))
print("val size",len(val_dataset))

```

```

train size 5000
test size 5000
val size 3000

```

```

[30]: def
    ↪ experiment(train_data, test_data, val_data, model_version, epochs, batchsize, learning_rate):
    ↪
        # initializing the train, validation, and test data loaders with given
    ↪ batchsize
        trainDataLoader = DataLoader(train_data, shuffle=True, batch_size=batchsize)
        valDataLoader = DataLoader(val_data, batch_size=batchsize)
        testDataLoader = DataLoader(test_data, batch_size=batchsize)

        # steps per epoch for train and validation set
        trainSteps = len(trainDataLoader.dataset) // batchsize
        valSteps = len(valDataLoader.dataset) // batchsize

        # initialize the model
        print("//...Initializing {} model...//".format(model_version))
        #model= LeNet5()
        model=model_version()

        #defining optimizer and cross-entropy loss function
        optimizer=Adam(model.parameters(), lr=learning_rate)
        criterion=nn.CrossEntropyLoss()

        #Learning rate will reduce by factor of 50% after every step size of
    ↪ 20epochs
        scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=20,
    ↪ gamma=0.5)

        #empty dict to store losses and accuracy score
        myDict={"train_loss": [], "train_acc": [], "val_loss": [], "val_acc": []}

        #initialize training the model
        print("//....Training initiated...//")
        start=time.time()

        tb = SummaryWriter()
        for epo in range(0, epochs):
            model.train()
            trainLoss=0
            trainCorrect=0
            valLoss=0
            valCorrect=0
            for (x,y) in trainDataLoader:
                optimizer.zero_grad()
                ytrain_pred=model(x)
                loss=criterion(ytrain_pred,y)
                loss.backward()

```

```

        optimizer.step()
        trainLoss+=loss
        correct_train=(ytrain_pred.argmax(1)==y).type(torch.float).sum()
    ↪item()

    trainCorrect=correct_train+trainCorrect
    with torch.no_grad():
        model.eval()
        for (x,y) in valDataLoader:
            yval_pred=model(x)
            loss=criterion(yval_pred,y)
            valLoss+=loss
            correct_val=(yval_pred.argmax(1)==y).type(torch.float).sum()
    ↪item()

        valCorrect=correct_val+valCorrect
    scheduler.step()

    #calculating loss and accuracy for each epoch for both train and val
    ↪sets

    trainLoss_avg=trainLoss/trainSteps
    valLoss_avg=valLoss/valSteps
    train_acc=trainCorrect/(len(trainDataLoader.dataset))
    val_acc=valCorrect/(len(valDataLoader.dataset))

    tb.add_scalar("train/Loss", trainLoss, epo)
    tb.add_scalar("train/Correct", trainCorrect, epo)
    tb.add_scalar("train/Accuracy", train_acc, epo)

    tb.add_scalar("val/Loss", valLoss, epo)
    tb.add_scalar("val/Correct", valCorrect, epo)
    tb.add_scalar("val/Accuracy", val_acc, epo)

    myDict["train_loss"].append(trainLoss_avg.cpu().detach().numpy())
    myDict["val_loss"].append(valLoss_avg.cpu().detach().numpy())
    myDict["train_acc"].append(train_acc)
    myDict["val_acc"].append(val_acc)

    print("//... epoch: {}/{}".format(epo + 1, epochs))
    print("Training loss: {:.4f}, Train accuracy: {:.2f}".
    ↪format(trainLoss_avg, train_acc))
    print("Validation loss: {:.4f}, Validation accuracy: {:.2f}\n".
    ↪format(valLoss_avg, val_acc))

    end=time.time()
    print("//...training the network took {:.2f} sec...//".format(end-start))

    #Evaluation of trained model on test set
    with torch.no_grad():

```

```

model.eval()
ytest_pred=[]
ytrue=[]
for (x,y) in testDataLoader:
    y_pred=model(x)
    ytest_pred.extend(y_pred.argmax(axis=1).numpy())
    ytrue.extend(y.numpy())

#print("y_pred",ytest_pred)
#print("y_true",ytrue)

class_names = train_data.classes
label_name = {'1': 'airplane', '2': 'bird', '3': 'car', '4': 'cat', '5':
↪ 'deer', '6': 'dog', '7': 'horse', '8': 'monkey', '9': 'ship', '10': 'truck'}
target_label=[]
for i in class_names:
    target_label.append(label_name[i])
#print("class_names",class_names)
#print("target_label",target_label)

print(classification_report(ytrue,ytest_pred,target_names=target_label))
print("//...Confusion Matrix...//")
print(confusion_matrix(ytrue, ytest_pred, labels=[x for x in range(10)]))

#plotting loss and acc
plt.style.use("ggplot")
plt.figure(2)
plt.plot(myDict["train_loss"], label="train_loss")
plt.plot(myDict["val_loss"], label="val_loss")
plt.plot(myDict["train_acc"], label="train_acc")
plt.plot(myDict["val_acc"], label="val_acc")
plt.title("Training Loss and Accuracy: {} model".format(type(model).
↪ __name__))
plt.xlabel("# of Epochs")
plt.ylabel("Loss/Accuracy")
plt.legend()
plt.close()

```

```

[39]: def ↵
↪ experiment_L2(train_data,test_data,val_data,model_version,epochs,batchsize,learning_rate):
↪
    # initializing the train, validation, and test data loaders with given ↵
↪ batchsize
    trainDataLoader = DataLoader(train_data, shuffle=True,batch_size=batchsize)
    valDataLoader = DataLoader(val_data, batch_size=batchsize)
    testDataLoader = DataLoader(test_data, batch_size=batchsize)

```



```

# steps per epoch for train and validation set
trainSteps = len(trainDataLoader.dataset) // batchsize
valSteps = len(valDataLoader.dataset) // batchsize

#initialize the model
print("//...Initializing {} model...//".format(model_version))
#model= LeNet5()
model=model_version()

#defining optimizer and cross-entropy loss function. Also Adding L2 loss by
→ specifying weight_decay
optimizer=Adam(model.parameters(),lr=learning_rate,weight_decay=1e-5)
criterion=nn.CrossEntropyLoss()

#Learning rate will reduce by factor of 50% after every step size of
→ 20epochs
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=20,
→ gamma=0.5)

#empty dict to store losses and accuracy score
myDict={"train_loss": [], "train_acc": [], "val_loss": [], "val_acc": []}

#initialize training the model
print("//....Training initiated...//")
start=time.time()

tb = SummaryWriter()
for epo in range(0,epochs):
    model.train()
    trainLoss=0
    trainCorrect=0
    valLoss=0
    valCorrect=0
    for (x,y) in trainDataLoader:
        optimizer.zero_grad()
        ytrain_pred=model(x)
        loss=criterion(ytrain_pred,y)
        loss.backward()
        optimizer.step()
        trainLoss+=loss
        correct_train=(ytrain_pred.argmax(1)==y).type(torch.float).sum().
→ item()
        trainCorrect=correct_train+trainCorrect
    with torch.no_grad():
        model.eval()
        for (x,y) in valDataLoader:

```

```

        yval_pred=model(x)
        loss=criterion(yval_pred,y)
        valLoss+=loss
        correct_val=(yval_pred.argmax(1)==y).type(torch.float).sum().
→item()

        valCorrect=correct_val+valCorrect
        scheduler.step()

#calculating loss and accuracy for each epoch for both train and val
→sets

        trainLoss_avg=trainLoss/trainSteps
        valLoss_avg=valLoss/valSteps
        train_acc=trainCorrect/(len(trainDataLoader.dataset))
        val_acc=valCorrect/(len(valDataLoader.dataset))

        tb.add_scalar("train/Loss", trainLoss, epo)
        tb.add_scalar("train/Correct", trainCorrect, epo)
        tb.add_scalar("train/Accuracy", train_acc, epo)

        tb.add_scalar("val/Loss", valLoss, epo)
        tb.add_scalar("val/Correct", valCorrect, epo)
        tb.add_scalar("val/Accuracy", val_acc, epo)

        myDict["train_loss"].append(trainLoss_avg.cpu().detach().numpy())
        myDict["val_loss"].append(valLoss_avg.cpu().detach().numpy())
        myDict["train_acc"].append(train_acc)
        myDict["val_acc"].append(val_acc)

        print("//... epoch: {}/{}".format(epo + 1, epochs))
        print("Training loss: {:.4f}, Train accuracy: {:.2f}".
→format(trainLoss_avg, train_acc))
        print("Validation loss: {:.4f}, Validation accuracy: {:.2f}\n".
→format(valLoss_avg, val_acc))

        end=time.time()
        print("//...training the network took {:.2f} sec...//".format(end-start))

#Evaluation of trained model on test set
        with torch.no_grad():
            model.eval()
            ytest_pred=[]
            ytrue=[]
            for (x,y) in testDataLoader:
                y_pred=model(x)
                ytest_pred.extend(y_pred.argmax(axis=1).numpy())
                ytrue.extend(y.numpy())

```

```

# print("y_pred", ytest_pred)
# print("y_true", ytrue)

class_names = train_data.classes
label_name = {'1': 'airplane', '2': 'bird', '3': 'car', '4': 'cat', '5':
↪ 'deer', '6': 'dog', '7': 'horse', '8': 'monkey', '9': 'ship', '10': 'truck'}
target_label = []
for i in class_names:
    target_label.append(label_name[i])
# print("class_names", class_names)
# print("target_label", target_label)

print(classification_report(ytrue, ytest_pred, target_names=target_label))
print("//...Confusion Matrix...//")
print(confusion_matrix(ytrue, ytest_pred, labels=[x for x in range(10)]))

# plotting loss and acc
plt.style.use("ggplot")
plt.figure(2)
plt.plot(myDict["train_loss"], label="train_loss")
plt.plot(myDict["val_loss"], label="val_loss")
plt.plot(myDict["train_acc"], label="train_acc")
plt.plot(myDict["val_acc"], label="val_acc")
plt.title("Training Loss and Accuracy: {} model".format(type(model).
↪ __name__))
plt.xlabel("# of Epochs")
plt.ylabel("Loss/Accuracy")
plt.legend()
plt.close()

```

```
[34]: experiment(train_dataset, test_dataset, val_dataset, LeNet5, 100, 128, 0.001)
```

```

//...Initializing <class '__main__.LeNet5'> model...//
//...Training initiated...//
//... epoch: 1/100
Training loss: 2.1893, Train accuracy: 0.22
Validation loss: 1.9880, Validation accuracy: 0.30

//... epoch: 2/100
Training loss: 1.9029, Train accuracy: 0.31
Validation loss: 1.8018, Validation accuracy: 0.35

//... epoch: 3/100
Training loss: 1.7870, Train accuracy: 0.35
Validation loss: 1.7513, Validation accuracy: 0.37

//... epoch: 4/100
Training loss: 1.7053, Train accuracy: 0.38

```

Validation loss: 1.6910, Validation accuracy: 0.38

//... epoch: 5/100

Training loss: 1.6349, Train accuracy: 0.40

Validation loss: 1.6592, Validation accuracy: 0.38

//... epoch: 6/100

Training loss: 1.5930, Train accuracy: 0.42

Validation loss: 1.5974, Validation accuracy: 0.42

//... epoch: 7/100

Training loss: 1.5472, Train accuracy: 0.43

Validation loss: 1.5909, Validation accuracy: 0.42

//... epoch: 8/100

Training loss: 1.4899, Train accuracy: 0.44

Validation loss: 1.5511, Validation accuracy: 0.44

//... epoch: 9/100

Training loss: 1.4532, Train accuracy: 0.49

Validation loss: 1.5359, Validation accuracy: 0.45

//... epoch: 10/100

Training loss: 1.4136, Train accuracy: 0.48

Validation loss: 1.5211, Validation accuracy: 0.45

//... epoch: 11/100

Training loss: 1.3705, Train accuracy: 0.50

Validation loss: 1.4971, Validation accuracy: 0.45

//... epoch: 12/100

Training loss: 1.3319, Train accuracy: 0.53

Validation loss: 1.4775, Validation accuracy: 0.47

//... epoch: 13/100

Training loss: 1.2696, Train accuracy: 0.54

Validation loss: 1.5481, Validation accuracy: 0.45

//... epoch: 14/100

Training loss: 1.2559, Train accuracy: 0.55

Validation loss: 1.4819, Validation accuracy: 0.47

//... epoch: 15/100

Training loss: 1.2275, Train accuracy: 0.56

Validation loss: 1.5157, Validation accuracy: 0.47

//... epoch: 16/100

Training loss: 1.2112, Train accuracy: 0.57

Validation loss: 1.4603, Validation accuracy: 0.49

//... epoch: 17/100

Training loss: 1.1306, Train accuracy: 0.60

Validation loss: 1.4551, Validation accuracy: 0.49

//... epoch: 18/100

Training loss: 1.0986, Train accuracy: 0.61

Validation loss: 1.5240, Validation accuracy: 0.47

//... epoch: 19/100

Training loss: 1.1025, Train accuracy: 0.62

Validation loss: 1.5698, Validation accuracy: 0.46

//... epoch: 20/100

Training loss: 1.1226, Train accuracy: 0.60

Validation loss: 1.5016, Validation accuracy: 0.48

//... epoch: 21/100

Training loss: 0.9601, Train accuracy: 0.66

Validation loss: 1.5028, Validation accuracy: 0.49

//... epoch: 22/100

Training loss: 0.9367, Train accuracy: 0.68

Validation loss: 1.5055, Validation accuracy: 0.50

//... epoch: 23/100

Training loss: 0.9085, Train accuracy: 0.69

Validation loss: 1.5155, Validation accuracy: 0.49

//... epoch: 24/100

Training loss: 0.8792, Train accuracy: 0.70

Validation loss: 1.5492, Validation accuracy: 0.50

//... epoch: 25/100

Training loss: 0.8604, Train accuracy: 0.71

Validation loss: 1.5350, Validation accuracy: 0.50

//... epoch: 26/100

Training loss: 0.8242, Train accuracy: 0.71

Validation loss: 1.5570, Validation accuracy: 0.49

//... epoch: 27/100

Training loss: 0.8049, Train accuracy: 0.72

Validation loss: 1.6138, Validation accuracy: 0.49

//... epoch: 28/100

Training loss: 0.7918, Train accuracy: 0.73

Validation loss: 1.6222, Validation accuracy: 0.48

//... epoch: 29/100

Training loss: 0.7628, Train accuracy: 0.74

Validation loss: 1.6111, Validation accuracy: 0.49

//... epoch: 30/100

Training loss: 0.7312, Train accuracy: 0.75

Validation loss: 1.6759, Validation accuracy: 0.49

//... epoch: 31/100

Training loss: 0.7362, Train accuracy: 0.75

Validation loss: 1.6729, Validation accuracy: 0.49

//... epoch: 32/100

Training loss: 0.6854, Train accuracy: 0.77

Validation loss: 1.6935, Validation accuracy: 0.49

//... epoch: 33/100

Training loss: 0.6496, Train accuracy: 0.78

Validation loss: 1.7520, Validation accuracy: 0.49

//... epoch: 34/100

Training loss: 0.6538, Train accuracy: 0.79

Validation loss: 1.7642, Validation accuracy: 0.49

//... epoch: 35/100

Training loss: 0.6180, Train accuracy: 0.79

Validation loss: 1.7563, Validation accuracy: 0.49

//... epoch: 36/100

Training loss: 0.5800, Train accuracy: 0.81

Validation loss: 1.8331, Validation accuracy: 0.49

//... epoch: 37/100

Training loss: 0.5553, Train accuracy: 0.82

Validation loss: 1.8434, Validation accuracy: 0.49

//... epoch: 38/100

Training loss: 0.5415, Train accuracy: 0.83

Validation loss: 1.9122, Validation accuracy: 0.49

//... epoch: 39/100

Training loss: 0.5334, Train accuracy: 0.83

Validation loss: 1.9637, Validation accuracy: 0.49

//... epoch: 40/100

Training loss: 0.4907, Train accuracy: 0.84

Validation loss: 1.9797, Validation accuracy: 0.48

//... epoch: 41/100

Training loss: 0.4518, Train accuracy: 0.86

Validation loss: 1.9894, Validation accuracy: 0.48

//... epoch: 42/100

Training loss: 0.4427, Train accuracy: 0.86

Validation loss: 2.0035, Validation accuracy: 0.48

//... epoch: 43/100

Training loss: 0.4200, Train accuracy: 0.88

Validation loss: 2.0682, Validation accuracy: 0.48

//... epoch: 44/100

Training loss: 0.4019, Train accuracy: 0.88

Validation loss: 2.0815, Validation accuracy: 0.48

//... epoch: 45/100

Training loss: 0.3942, Train accuracy: 0.89

Validation loss: 2.1038, Validation accuracy: 0.48

//... epoch: 46/100

Training loss: 0.3817, Train accuracy: 0.89

Validation loss: 2.1252, Validation accuracy: 0.48

//... epoch: 47/100

Training loss: 0.3768, Train accuracy: 0.89

Validation loss: 2.1500, Validation accuracy: 0.48

//... epoch: 48/100

Training loss: 0.3652, Train accuracy: 0.89

Validation loss: 2.1895, Validation accuracy: 0.48

//... epoch: 49/100

Training loss: 0.3408, Train accuracy: 0.91

Validation loss: 2.2270, Validation accuracy: 0.48

//... epoch: 50/100

Training loss: 0.3400, Train accuracy: 0.91

Validation loss: 2.2517, Validation accuracy: 0.48

//... epoch: 51/100

Training loss: 0.3448, Train accuracy: 0.90

Validation loss: 2.2656, Validation accuracy: 0.47

//... epoch: 52/100

Training loss: 0.3339, Train accuracy: 0.91

Validation loss: 2.2848, Validation accuracy: 0.48

//... epoch: 53/100

Training loss: 0.3075, Train accuracy: 0.92

Validation loss: 2.3376, Validation accuracy: 0.47

//... epoch: 54/100

Training loss: 0.2970, Train accuracy: 0.92

Validation loss: 2.3795, Validation accuracy: 0.48

//... epoch: 55/100

Training loss: 0.3021, Train accuracy: 0.92

Validation loss: 2.4018, Validation accuracy: 0.47

//... epoch: 56/100

Training loss: 0.2858, Train accuracy: 0.93

Validation loss: 2.4166, Validation accuracy: 0.47

//... epoch: 57/100

Training loss: 0.2754, Train accuracy: 0.93

Validation loss: 2.4823, Validation accuracy: 0.47

//... epoch: 58/100

Training loss: 0.2837, Train accuracy: 0.93

Validation loss: 2.4877, Validation accuracy: 0.48

//... epoch: 59/100

Training loss: 0.2685, Train accuracy: 0.93

Validation loss: 2.5258, Validation accuracy: 0.47

//... epoch: 60/100

Training loss: 0.2455, Train accuracy: 0.94

Validation loss: 2.5445, Validation accuracy: 0.47

//... epoch: 61/100

Training loss: 0.2314, Train accuracy: 0.95

Validation loss: 2.5834, Validation accuracy: 0.47

//... epoch: 62/100

Training loss: 0.2190, Train accuracy: 0.95

Validation loss: 2.6038, Validation accuracy: 0.47

//... epoch: 63/100

Training loss: 0.2181, Train accuracy: 0.95

Validation loss: 2.6277, Validation accuracy: 0.47

//... epoch: 64/100

Training loss: 0.2124, Train accuracy: 0.95

Validation loss: 2.6448, Validation accuracy: 0.47

//... epoch: 65/100

Training loss: 0.2171, Train accuracy: 0.95

Validation loss: 2.6649, Validation accuracy: 0.47

//... epoch: 66/100

Training loss: 0.2161, Train accuracy: 0.95

Validation loss: 2.6741, Validation accuracy: 0.46

//... epoch: 67/100

Training loss: 0.2028, Train accuracy: 0.96

Validation loss: 2.6877, Validation accuracy: 0.47

//... epoch: 68/100

Training loss: 0.1943, Train accuracy: 0.96

Validation loss: 2.7245, Validation accuracy: 0.47

//... epoch: 69/100

Training loss: 0.1906, Train accuracy: 0.96

Validation loss: 2.7265, Validation accuracy: 0.46

//... epoch: 70/100

Training loss: 0.1875, Train accuracy: 0.96

Validation loss: 2.7588, Validation accuracy: 0.47

//... epoch: 71/100

Training loss: 0.1843, Train accuracy: 0.96

Validation loss: 2.7839, Validation accuracy: 0.46

//... epoch: 72/100

Training loss: 0.1871, Train accuracy: 0.96

Validation loss: 2.8084, Validation accuracy: 0.46

//... epoch: 73/100

Training loss: 0.1838, Train accuracy: 0.96

Validation loss: 2.8208, Validation accuracy: 0.46

//... epoch: 74/100

Training loss: 0.1744, Train accuracy: 0.97

Validation loss: 2.8378, Validation accuracy: 0.46

//... epoch: 75/100

Training loss: 0.1725, Train accuracy: 0.96

Validation loss: 2.8574, Validation accuracy: 0.47

//... epoch: 76/100

Training loss: 0.1699, Train accuracy: 0.97

Validation loss: 2.8930, Validation accuracy: 0.46

//... epoch: 77/100

Training loss: 0.1656, Train accuracy: 0.97

Validation loss: 2.9076, Validation accuracy: 0.46

//... epoch: 78/100

Training loss: 0.1624, Train accuracy: 0.97

Validation loss: 2.9186, Validation accuracy: 0.46

//... epoch: 79/100

Training loss: 0.1583, Train accuracy: 0.97

Validation loss: 2.9455, Validation accuracy: 0.46

//... epoch: 80/100

Training loss: 0.1536, Train accuracy: 0.97

Validation loss: 2.9672, Validation accuracy: 0.46

//... epoch: 81/100

Training loss: 0.1459, Train accuracy: 0.97

Validation loss: 2.9869, Validation accuracy: 0.46

//... epoch: 82/100

Training loss: 0.1426, Train accuracy: 0.97

Validation loss: 2.9919, Validation accuracy: 0.46

//... epoch: 83/100

Training loss: 0.1400, Train accuracy: 0.97

Validation loss: 3.0093, Validation accuracy: 0.46

//... epoch: 84/100

Training loss: 0.1431, Train accuracy: 0.98

Validation loss: 3.0260, Validation accuracy: 0.46

//... epoch: 85/100

Training loss: 0.1371, Train accuracy: 0.98

Validation loss: 3.0300, Validation accuracy: 0.46

//... epoch: 86/100

Training loss: 0.1349, Train accuracy: 0.98

Validation loss: 3.0438, Validation accuracy: 0.46

//... epoch: 87/100

Training loss: 0.1324, Train accuracy: 0.98

Validation loss: 3.0618, Validation accuracy: 0.46

//... epoch: 88/100

Training loss: 0.1341, Train accuracy: 0.98

Validation loss: 3.0713, Validation accuracy: 0.46

//... epoch: 89/100

Training loss: 0.1341, Train accuracy: 0.98

Validation loss: 3.0813, Validation accuracy: 0.46

//... epoch: 90/100

Training loss: 0.1306, Train accuracy: 0.98

Validation loss: 3.0864, Validation accuracy: 0.46

//... epoch: 91/100

Training loss: 0.1262, Train accuracy: 0.98

Validation loss: 3.1080, Validation accuracy: 0.46

//... epoch: 92/100

Training loss: 0.1241, Train accuracy: 0.98

Validation loss: 3.1230, Validation accuracy: 0.46

//... epoch: 93/100

Training loss: 0.1251, Train accuracy: 0.98

Validation loss: 3.1401, Validation accuracy: 0.46

//... epoch: 94/100

Training loss: 0.1305, Train accuracy: 0.98

Validation loss: 3.1447, Validation accuracy: 0.46

//... epoch: 95/100

Training loss: 0.1246, Train accuracy: 0.98

Validation loss: 3.1621, Validation accuracy: 0.46

//... epoch: 96/100

Training loss: 0.1207, Train accuracy: 0.98

Validation loss: 3.1814, Validation accuracy: 0.46

//... epoch: 97/100

Training loss: 0.1268, Train accuracy: 0.98

Validation loss: 3.1911, Validation accuracy: 0.46

//... epoch: 98/100

Training loss: 0.1239, Train accuracy: 0.98

Validation loss: 3.2100, Validation accuracy: 0.46

//... epoch: 99/100

Training loss: 0.1148, Train accuracy: 0.98

Validation loss: 3.2228, Validation accuracy: 0.46

//... epoch: 100/100

Training loss: 0.1184, Train accuracy: 0.98

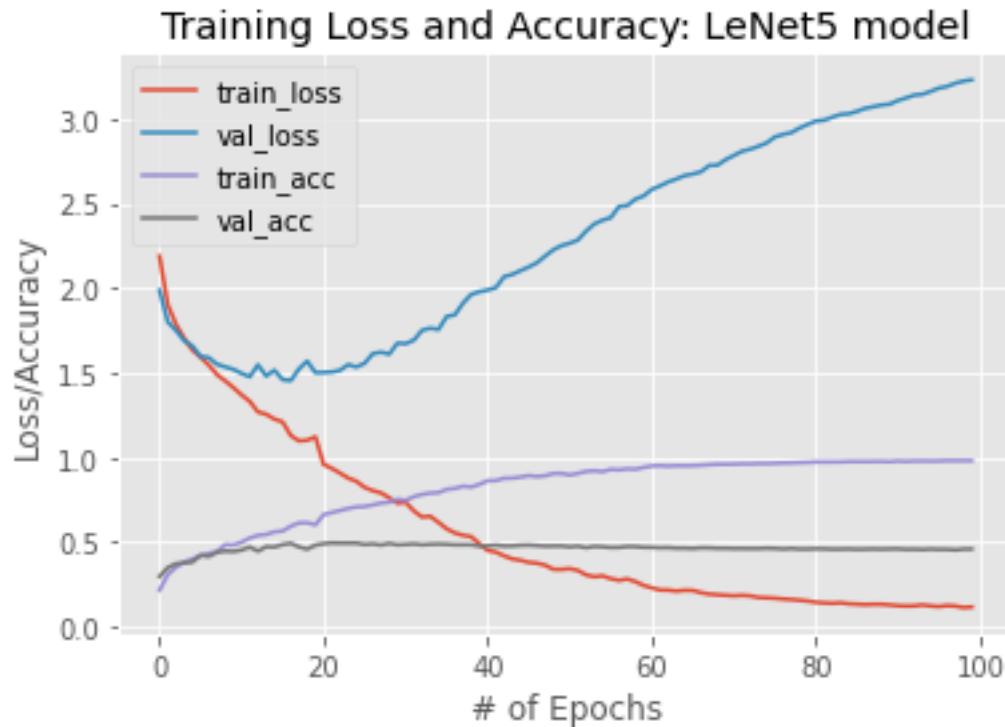
Validation loss: 3.2306, Validation accuracy: 0.46

//...training the network took 879.43 sec...//

	precision	recall	f1-score	support
airplane	0.64	0.65	0.64	500
truck	0.53	0.53	0.53	500
bird	0.35	0.33	0.34	500
car	0.64	0.55	0.59	500
cat	0.30	0.32	0.31	500
deer	0.40	0.42	0.41	500
dog	0.27	0.26	0.27	500
horse	0.50	0.48	0.49	500
monkey	0.33	0.36	0.34	500
ship	0.64	0.67	0.65	500
accuracy			0.46	5000
macro avg	0.46	0.46	0.46	5000
weighted avg	0.46	0.46	0.46	5000

//...Confusion Matrix...//

```
[[323 35 17 18 11 22 10 3 8 53]
 [ 35 264 12 78 10 8 13 9 12 59]
 [ 26 10 167 9 72 62 45 23 76 10]
 [ 24 96 16 273 17 9 8 6 8 43]
 [ 9 13 56 10 162 55 64 38 78 15]
 [ 12 7 57 4 56 212 48 48 53 3]
 [ 10 5 55 4 81 57 130 60 94 4]
 [ 4 14 30 3 33 50 85 240 40 1]
 [ 9 7 65 2 72 40 73 49 180 3]
 [ 56 43 5 24 18 12 4 1 4 333]]
```



```
[38]: experiment(train_dataset,test_dataset,val_dataset,LeNet5_BN,100,128,0.001)
```

```
//...Initializing <class '__main__.LeNet5_BN'> model...//
//...Training initiated...//
//... epoch: 1/100
Training loss: 2.0004, Train accuracy: 0.29
Validation loss: 1.8269, Validation accuracy: 0.38

//... epoch: 2/100
Training loss: 1.6341, Train accuracy: 0.43
Validation loss: 1.6214, Validation accuracy: 0.44

//... epoch: 3/100
Training loss: 1.4449, Train accuracy: 0.50
Validation loss: 1.5258, Validation accuracy: 0.44

//... epoch: 4/100
Training loss: 1.2963, Train accuracy: 0.55
Validation loss: 1.4364, Validation accuracy: 0.48

//... epoch: 5/100
Training loss: 1.1437, Train accuracy: 0.62
Validation loss: 1.4283, Validation accuracy: 0.48
```

//... epoch: 6/100
Training loss: 1.0576, Train accuracy: 0.65
Validation loss: 1.4734, Validation accuracy: 0.47

//... epoch: 7/100
Training loss: 0.9535, Train accuracy: 0.68
Validation loss: 1.4347, Validation accuracy: 0.50

//... epoch: 8/100
Training loss: 0.8492, Train accuracy: 0.74
Validation loss: 1.6364, Validation accuracy: 0.46

//... epoch: 9/100
Training loss: 0.8075, Train accuracy: 0.74
Validation loss: 1.5303, Validation accuracy: 0.48

//... epoch: 10/100
Training loss: 0.7378, Train accuracy: 0.77
Validation loss: 1.5850, Validation accuracy: 0.48

//... epoch: 11/100
Training loss: 0.6188, Train accuracy: 0.82
Validation loss: 1.5557, Validation accuracy: 0.49

//... epoch: 12/100
Training loss: 0.5384, Train accuracy: 0.84
Validation loss: 1.7014, Validation accuracy: 0.47

//... epoch: 13/100
Training loss: 0.5202, Train accuracy: 0.85
Validation loss: 1.7364, Validation accuracy: 0.46

//... epoch: 14/100
Training loss: 0.4248, Train accuracy: 0.88
Validation loss: 1.7711, Validation accuracy: 0.47

//... epoch: 15/100
Training loss: 0.3563, Train accuracy: 0.91
Validation loss: 1.8509, Validation accuracy: 0.47

//... epoch: 16/100
Training loss: 0.2912, Train accuracy: 0.94
Validation loss: 1.9018, Validation accuracy: 0.47

//... epoch: 17/100
Training loss: 0.2945, Train accuracy: 0.93
Validation loss: 1.9657, Validation accuracy: 0.47

//... epoch: 18/100
Training loss: 0.2489, Train accuracy: 0.95
Validation loss: 1.9951, Validation accuracy: 0.47

//... epoch: 19/100
Training loss: 0.3013, Train accuracy: 0.92
Validation loss: 2.0768, Validation accuracy: 0.47

//... epoch: 20/100
Training loss: 0.2279, Train accuracy: 0.94
Validation loss: 2.0927, Validation accuracy: 0.47

//... epoch: 21/100
Training loss: 0.1859, Train accuracy: 0.97
Validation loss: 2.0807, Validation accuracy: 0.47

//... epoch: 22/100
Training loss: 0.1787, Train accuracy: 0.96
Validation loss: 2.1159, Validation accuracy: 0.47

//... epoch: 23/100
Training loss: 0.1343, Train accuracy: 0.98
Validation loss: 2.1636, Validation accuracy: 0.47

//... epoch: 24/100
Training loss: 0.1090, Train accuracy: 0.99
Validation loss: 2.2064, Validation accuracy: 0.47

//... epoch: 25/100
Training loss: 0.1024, Train accuracy: 0.99
Validation loss: 2.2353, Validation accuracy: 0.47

//... epoch: 26/100
Training loss: 0.1211, Train accuracy: 0.99
Validation loss: 2.2390, Validation accuracy: 0.47

//... epoch: 27/100
Training loss: 0.1306, Train accuracy: 0.98
Validation loss: 2.3102, Validation accuracy: 0.47

//... epoch: 28/100
Training loss: 0.1263, Train accuracy: 0.98
Validation loss: 2.3004, Validation accuracy: 0.47

//... epoch: 29/100
Training loss: 0.1426, Train accuracy: 0.98
Validation loss: 2.3217, Validation accuracy: 0.46

//... epoch: 30/100
Training loss: 0.1165, Train accuracy: 0.98
Validation loss: 2.4132, Validation accuracy: 0.46

//... epoch: 31/100
Training loss: 0.1316, Train accuracy: 0.97
Validation loss: 2.3990, Validation accuracy: 0.46

//... epoch: 32/100
Training loss: 0.0909, Train accuracy: 0.99
Validation loss: 2.3791, Validation accuracy: 0.46

//... epoch: 33/100
Training loss: 0.0634, Train accuracy: 1.00
Validation loss: 2.4255, Validation accuracy: 0.46

//... epoch: 34/100
Training loss: 0.0757, Train accuracy: 0.99
Validation loss: 2.4597, Validation accuracy: 0.46

//... epoch: 35/100
Training loss: 0.1175, Train accuracy: 0.99
Validation loss: 2.4815, Validation accuracy: 0.47

//... epoch: 36/100
Training loss: 0.0988, Train accuracy: 0.98
Validation loss: 2.5337, Validation accuracy: 0.46

//... epoch: 37/100
Training loss: 0.0676, Train accuracy: 1.00
Validation loss: 2.5713, Validation accuracy: 0.46

//... epoch: 38/100
Training loss: 0.1229, Train accuracy: 0.98
Validation loss: 2.5859, Validation accuracy: 0.46

//... epoch: 39/100
Training loss: 0.1001, Train accuracy: 0.98
Validation loss: 2.5449, Validation accuracy: 0.47

//... epoch: 40/100
Training loss: 0.0849, Train accuracy: 0.99
Validation loss: 2.6349, Validation accuracy: 0.46

//... epoch: 41/100
Training loss: 0.0811, Train accuracy: 0.98
Validation loss: 2.5659, Validation accuracy: 0.46

//... epoch: 42/100
Training loss: 0.0562, Train accuracy: 1.00
Validation loss: 2.6020, Validation accuracy: 0.46

//... epoch: 43/100
Training loss: 0.0543, Train accuracy: 1.00
Validation loss: 2.5712, Validation accuracy: 0.48

//... epoch: 44/100
Training loss: 0.0609, Train accuracy: 1.00
Validation loss: 2.6131, Validation accuracy: 0.47

//... epoch: 45/100
Training loss: 0.0849, Train accuracy: 0.99
Validation loss: 2.6647, Validation accuracy: 0.46

//... epoch: 46/100
Training loss: 0.0645, Train accuracy: 0.99
Validation loss: 2.6266, Validation accuracy: 0.46

//... epoch: 47/100
Training loss: 0.0574, Train accuracy: 1.00
Validation loss: 2.6173, Validation accuracy: 0.47

//... epoch: 48/100
Training loss: 0.0490, Train accuracy: 1.00
Validation loss: 2.6104, Validation accuracy: 0.47

//... epoch: 49/100
Training loss: 0.0470, Train accuracy: 1.00
Validation loss: 2.6562, Validation accuracy: 0.47

//... epoch: 50/100
Training loss: 0.0608, Train accuracy: 0.99
Validation loss: 2.6214, Validation accuracy: 0.47

//... epoch: 51/100
Training loss: 0.0466, Train accuracy: 1.00
Validation loss: 2.6553, Validation accuracy: 0.47

//... epoch: 52/100
Training loss: 0.0435, Train accuracy: 1.00
Validation loss: 2.6477, Validation accuracy: 0.46

//... epoch: 53/100
Training loss: 0.0464, Train accuracy: 0.99
Validation loss: 2.6555, Validation accuracy: 0.47

//... epoch: 54/100
Training loss: 0.0301, Train accuracy: 1.00
Validation loss: 2.6291, Validation accuracy: 0.47

//... epoch: 55/100
Training loss: 0.0324, Train accuracy: 1.00
Validation loss: 2.6636, Validation accuracy: 0.47

//... epoch: 56/100
Training loss: 0.0518, Train accuracy: 1.00
Validation loss: 2.6636, Validation accuracy: 0.47

//... epoch: 57/100
Training loss: 0.0555, Train accuracy: 1.00
Validation loss: 2.7193, Validation accuracy: 0.47

//... epoch: 58/100
Training loss: 0.0523, Train accuracy: 0.99
Validation loss: 2.7290, Validation accuracy: 0.46

//... epoch: 59/100
Training loss: 0.0444, Train accuracy: 1.00
Validation loss: 2.7328, Validation accuracy: 0.47

//... epoch: 60/100
Training loss: 0.0535, Train accuracy: 0.99
Validation loss: 2.7316, Validation accuracy: 0.46

//... epoch: 61/100
Training loss: 0.0308, Train accuracy: 1.00
Validation loss: 2.7362, Validation accuracy: 0.46

//... epoch: 62/100
Training loss: 0.0434, Train accuracy: 1.00
Validation loss: 2.7505, Validation accuracy: 0.46

//... epoch: 63/100
Training loss: 0.0233, Train accuracy: 1.00
Validation loss: 2.7327, Validation accuracy: 0.46

//... epoch: 64/100
Training loss: 0.0395, Train accuracy: 1.00
Validation loss: 2.7456, Validation accuracy: 0.47

//... epoch: 65/100
Training loss: 0.0441, Train accuracy: 1.00
Validation loss: 2.7402, Validation accuracy: 0.47

//... epoch: 66/100
Training loss: 0.0469, Train accuracy: 1.00
Validation loss: 2.7773, Validation accuracy: 0.47

//... epoch: 67/100
Training loss: 0.0466, Train accuracy: 1.00
Validation loss: 2.7989, Validation accuracy: 0.46

//... epoch: 68/100
Training loss: 0.0602, Train accuracy: 1.00
Validation loss: 2.7723, Validation accuracy: 0.46

//... epoch: 69/100
Training loss: 0.0687, Train accuracy: 1.00
Validation loss: 2.8072, Validation accuracy: 0.47

//... epoch: 70/100
Training loss: 0.0490, Train accuracy: 1.00
Validation loss: 2.7774, Validation accuracy: 0.47

//... epoch: 71/100
Training loss: 0.0318, Train accuracy: 1.00
Validation loss: 2.7608, Validation accuracy: 0.47

//... epoch: 72/100
Training loss: 0.0318, Train accuracy: 1.00
Validation loss: 2.7886, Validation accuracy: 0.47

//... epoch: 73/100
Training loss: 0.0387, Train accuracy: 1.00
Validation loss: 2.7920, Validation accuracy: 0.47

//... epoch: 74/100
Training loss: 0.0413, Train accuracy: 1.00
Validation loss: 2.8145, Validation accuracy: 0.47

//... epoch: 75/100
Training loss: 0.0418, Train accuracy: 1.00
Validation loss: 2.7859, Validation accuracy: 0.47

//... epoch: 76/100
Training loss: 0.0438, Train accuracy: 1.00
Validation loss: 2.8010, Validation accuracy: 0.47

//... epoch: 77/100
Training loss: 0.0391, Train accuracy: 1.00
Validation loss: 2.7997, Validation accuracy: 0.47

//... epoch: 78/100
Training loss: 0.0281, Train accuracy: 1.00
Validation loss: 2.7846, Validation accuracy: 0.47

//... epoch: 79/100
Training loss: 0.0326, Train accuracy: 1.00
Validation loss: 2.7615, Validation accuracy: 0.47

//... epoch: 80/100
Training loss: 0.0337, Train accuracy: 1.00
Validation loss: 2.7913, Validation accuracy: 0.47

//... epoch: 81/100
Training loss: 0.0418, Train accuracy: 1.00
Validation loss: 2.8234, Validation accuracy: 0.47

//... epoch: 82/100
Training loss: 0.0398, Train accuracy: 1.00
Validation loss: 2.7577, Validation accuracy: 0.47

//... epoch: 83/100
Training loss: 0.0506, Train accuracy: 1.00
Validation loss: 2.7633, Validation accuracy: 0.47

//... epoch: 84/100
Training loss: 0.0299, Train accuracy: 1.00
Validation loss: 2.7671, Validation accuracy: 0.47

//... epoch: 85/100
Training loss: 0.0338, Train accuracy: 1.00
Validation loss: 2.7853, Validation accuracy: 0.47

//... epoch: 86/100
Training loss: 0.0258, Train accuracy: 1.00
Validation loss: 2.8067, Validation accuracy: 0.47

//... epoch: 87/100
Training loss: 0.0417, Train accuracy: 1.00
Validation loss: 2.8098, Validation accuracy: 0.47

//... epoch: 88/100
Training loss: 0.0486, Train accuracy: 1.00
Validation loss: 2.8239, Validation accuracy: 0.47

//... epoch: 89/100
Training loss: 0.0199, Train accuracy: 1.00
Validation loss: 2.8054, Validation accuracy: 0.47

```

//... epoch: 90/100
Training loss: 0.0317, Train accuracy: 1.00
Validation loss: 2.8123, Validation accuracy: 0.47

//... epoch: 91/100
Training loss: 0.0259, Train accuracy: 1.00
Validation loss: 2.8150, Validation accuracy: 0.47

//... epoch: 92/100
Training loss: 0.0406, Train accuracy: 1.00
Validation loss: 2.8277, Validation accuracy: 0.47

//... epoch: 93/100
Training loss: 0.0160, Train accuracy: 1.00
Validation loss: 2.7902, Validation accuracy: 0.46

//... epoch: 94/100
Training loss: 0.0222, Train accuracy: 1.00
Validation loss: 2.8152, Validation accuracy: 0.47

//... epoch: 95/100
Training loss: 0.0293, Train accuracy: 1.00
Validation loss: 2.7899, Validation accuracy: 0.47

//... epoch: 96/100
Training loss: 0.0173, Train accuracy: 1.00
Validation loss: 2.7930, Validation accuracy: 0.47

//... epoch: 97/100
Training loss: 0.0152, Train accuracy: 1.00
Validation loss: 2.7906, Validation accuracy: 0.48

//... epoch: 98/100
Training loss: 0.0332, Train accuracy: 1.00
Validation loss: 2.7950, Validation accuracy: 0.47

//... epoch: 99/100
Training loss: 0.0272, Train accuracy: 1.00
Validation loss: 2.7816, Validation accuracy: 0.47

//... epoch: 100/100
Training loss: 0.0178, Train accuracy: 1.00
Validation loss: 2.8399, Validation accuracy: 0.47

//...training the network took 919.00 sec...//
      precision    recall  f1-score   support


airplane      0.69      0.71      0.70        500

```

truck	0.53	0.45	0.49	500
bird	0.38	0.31	0.34	500
car	0.67	0.63	0.65	500
cat	0.28	0.28	0.28	500
deer	0.40	0.42	0.41	500
dog	0.26	0.29	0.28	500
horse	0.45	0.51	0.48	500
monkey	0.33	0.34	0.33	500
ship	0.64	0.68	0.66	500
accuracy			0.46	5000
macro avg	0.46	0.46	0.46	5000
weighted avg	0.46	0.46	0.46	5000

//...Confusion Matrix...//

```
[[353  25  15  17   6   9  10  10   5  50]
 [ 28 224  11  86  23  20  15  16   7  70]
 [ 24   5 154   7  60  63  66  31  73  17]
 [ 16  73  18 313  18   4  13  10  11  24]
 [  9  11  42  15 141  82  72  44  72  12]
 [ 11   5  46   4  65 212  40  60  50   7]
 [  7  10  49   3  74  55 147  82  68   5]
 [  4  14  30   4  33  31  83 255  44   2]
 [  6   6  27   3  70  49 112  56 168   3]
 [ 52  46   9  12  13   8   6   7   8 339]]
```



```
[42]: experiment_L2(train_dataset,test_dataset,val_dataset,LeNet5,100,128,0.001)
```

```
//...Initializing <class '__main__.LeNet5'> model...//
//...Training initiated...//
//... epoch: 1/100
Training loss: 2.1840, Train accuracy: 0.21
Validation loss: 2.0263, Validation accuracy: 0.28

//... epoch: 2/100
Training loss: 1.8906, Train accuracy: 0.31
Validation loss: 1.8038, Validation accuracy: 0.35

//... epoch: 3/100
Training loss: 1.7384, Train accuracy: 0.37
Validation loss: 1.7840, Validation accuracy: 0.36

//... epoch: 4/100
Training loss: 1.6290, Train accuracy: 0.40
Validation loss: 1.6395, Validation accuracy: 0.40

//... epoch: 5/100
Training loss: 1.5634, Train accuracy: 0.43
Validation loss: 1.6028, Validation accuracy: 0.41

//... epoch: 6/100
Training loss: 1.5148, Train accuracy: 0.45
Validation loss: 1.5749, Validation accuracy: 0.41

//... epoch: 7/100
Training loss: 1.4913, Train accuracy: 0.45
Validation loss: 1.5488, Validation accuracy: 0.44

//... epoch: 8/100
Training loss: 1.4514, Train accuracy: 0.46
Validation loss: 1.5401, Validation accuracy: 0.44

//... epoch: 9/100
Training loss: 1.4151, Train accuracy: 0.48
Validation loss: 1.5380, Validation accuracy: 0.44

//... epoch: 10/100
Training loss: 1.3781, Train accuracy: 0.51
Validation loss: 1.5531, Validation accuracy: 0.43

//... epoch: 11/100
```

Training loss: 1.3777, Train accuracy: 0.50
Validation loss: 1.4995, Validation accuracy: 0.45

//... epoch: 12/100
Training loss: 1.3007, Train accuracy: 0.53
Validation loss: 1.5467, Validation accuracy: 0.46

//... epoch: 13/100
Training loss: 1.3060, Train accuracy: 0.53
Validation loss: 1.5058, Validation accuracy: 0.45

//... epoch: 14/100
Training loss: 1.2497, Train accuracy: 0.54
Validation loss: 1.5459, Validation accuracy: 0.46

//... epoch: 15/100
Training loss: 1.2066, Train accuracy: 0.55
Validation loss: 1.5125, Validation accuracy: 0.47

//... epoch: 16/100
Training loss: 1.1822, Train accuracy: 0.58
Validation loss: 1.5244, Validation accuracy: 0.48

//... epoch: 17/100
Training loss: 1.1506, Train accuracy: 0.59
Validation loss: 1.5289, Validation accuracy: 0.47

//... epoch: 18/100
Training loss: 1.1144, Train accuracy: 0.61
Validation loss: 1.4936, Validation accuracy: 0.49

//... epoch: 19/100
Training loss: 1.0703, Train accuracy: 0.61
Validation loss: 1.4977, Validation accuracy: 0.48

//... epoch: 20/100
Training loss: 1.0458, Train accuracy: 0.63
Validation loss: 1.5291, Validation accuracy: 0.47

//... epoch: 21/100
Training loss: 0.9673, Train accuracy: 0.66
Validation loss: 1.5597, Validation accuracy: 0.48

//... epoch: 22/100
Training loss: 0.9562, Train accuracy: 0.66
Validation loss: 1.5317, Validation accuracy: 0.48

//... epoch: 23/100

Training loss: 0.9183, Train accuracy: 0.68
Validation loss: 1.5437, Validation accuracy: 0.48

//... epoch: 24/100
Training loss: 0.8843, Train accuracy: 0.68
Validation loss: 1.5576, Validation accuracy: 0.49

//... epoch: 25/100
Training loss: 0.8693, Train accuracy: 0.70
Validation loss: 1.5598, Validation accuracy: 0.48

//... epoch: 26/100
Training loss: 0.8598, Train accuracy: 0.71
Validation loss: 1.6038, Validation accuracy: 0.48

//... epoch: 27/100
Training loss: 0.8292, Train accuracy: 0.71
Validation loss: 1.6029, Validation accuracy: 0.48

//... epoch: 28/100
Training loss: 0.8310, Train accuracy: 0.72
Validation loss: 1.6075, Validation accuracy: 0.48

//... epoch: 29/100
Training loss: 0.7866, Train accuracy: 0.73
Validation loss: 1.6268, Validation accuracy: 0.48

//... epoch: 30/100
Training loss: 0.7474, Train accuracy: 0.74
Validation loss: 1.6721, Validation accuracy: 0.48

//... epoch: 31/100
Training loss: 0.7692, Train accuracy: 0.74
Validation loss: 1.6731, Validation accuracy: 0.48

//... epoch: 32/100
Training loss: 0.7193, Train accuracy: 0.75
Validation loss: 1.7210, Validation accuracy: 0.48

//... epoch: 33/100
Training loss: 0.7156, Train accuracy: 0.76
Validation loss: 1.6966, Validation accuracy: 0.48

//... epoch: 34/100
Training loss: 0.6906, Train accuracy: 0.76
Validation loss: 1.7905, Validation accuracy: 0.48

//... epoch: 35/100

Training loss: 0.6743, Train accuracy: 0.77
Validation loss: 1.7951, Validation accuracy: 0.47

//... epoch: 36/100
Training loss: 0.6257, Train accuracy: 0.79
Validation loss: 1.7878, Validation accuracy: 0.48

//... epoch: 37/100
Training loss: 0.6168, Train accuracy: 0.79
Validation loss: 1.8189, Validation accuracy: 0.47

//... epoch: 38/100
Training loss: 0.6073, Train accuracy: 0.80
Validation loss: 1.8444, Validation accuracy: 0.48

//... epoch: 39/100
Training loss: 0.5825, Train accuracy: 0.80
Validation loss: 1.8840, Validation accuracy: 0.48

//... epoch: 40/100
Training loss: 0.5741, Train accuracy: 0.81
Validation loss: 1.9016, Validation accuracy: 0.48

//... epoch: 41/100
Training loss: 0.5195, Train accuracy: 0.83
Validation loss: 1.9121, Validation accuracy: 0.48

//... epoch: 42/100
Training loss: 0.4959, Train accuracy: 0.84
Validation loss: 1.9309, Validation accuracy: 0.48

//... epoch: 43/100
Training loss: 0.4945, Train accuracy: 0.85
Validation loss: 1.9457, Validation accuracy: 0.47

//... epoch: 44/100
Training loss: 0.4820, Train accuracy: 0.85
Validation loss: 1.9631, Validation accuracy: 0.48

//... epoch: 45/100
Training loss: 0.4655, Train accuracy: 0.85
Validation loss: 2.0087, Validation accuracy: 0.48

//... epoch: 46/100
Training loss: 0.4502, Train accuracy: 0.86
Validation loss: 2.0580, Validation accuracy: 0.48

//... epoch: 47/100

Training loss: 0.4613, Train accuracy: 0.86
Validation loss: 2.0415, Validation accuracy: 0.47

//... epoch: 48/100
Training loss: 0.4410, Train accuracy: 0.87
Validation loss: 2.0853, Validation accuracy: 0.47

//... epoch: 49/100
Training loss: 0.4307, Train accuracy: 0.87
Validation loss: 2.0808, Validation accuracy: 0.48

//... epoch: 50/100
Training loss: 0.4149, Train accuracy: 0.87
Validation loss: 2.1275, Validation accuracy: 0.48

//... epoch: 51/100
Training loss: 0.4072, Train accuracy: 0.88
Validation loss: 2.1489, Validation accuracy: 0.47

//... epoch: 52/100
Training loss: 0.3978, Train accuracy: 0.88
Validation loss: 2.1600, Validation accuracy: 0.48

//... epoch: 53/100
Training loss: 0.3877, Train accuracy: 0.89
Validation loss: 2.1768, Validation accuracy: 0.48

//... epoch: 54/100
Training loss: 0.3852, Train accuracy: 0.88
Validation loss: 2.1956, Validation accuracy: 0.48

//... epoch: 55/100
Training loss: 0.3642, Train accuracy: 0.89
Validation loss: 2.2482, Validation accuracy: 0.48

//... epoch: 56/100
Training loss: 0.3564, Train accuracy: 0.89
Validation loss: 2.2775, Validation accuracy: 0.47

//... epoch: 57/100
Training loss: 0.3534, Train accuracy: 0.89
Validation loss: 2.2768, Validation accuracy: 0.48

//... epoch: 58/100
Training loss: 0.3403, Train accuracy: 0.90
Validation loss: 2.3356, Validation accuracy: 0.47

//... epoch: 59/100

Training loss: 0.3404, Train accuracy: 0.90
Validation loss: 2.3579, Validation accuracy: 0.48

//... epoch: 60/100
Training loss: 0.3206, Train accuracy: 0.91
Validation loss: 2.3681, Validation accuracy: 0.47

//... epoch: 61/100
Training loss: 0.2990, Train accuracy: 0.92
Validation loss: 2.3852, Validation accuracy: 0.48

//... epoch: 62/100
Training loss: 0.2980, Train accuracy: 0.92
Validation loss: 2.3974, Validation accuracy: 0.48

//... epoch: 63/100
Training loss: 0.2837, Train accuracy: 0.92
Validation loss: 2.4253, Validation accuracy: 0.47

//... epoch: 64/100
Training loss: 0.2827, Train accuracy: 0.93
Validation loss: 2.4338, Validation accuracy: 0.47

//... epoch: 65/100
Training loss: 0.2792, Train accuracy: 0.93
Validation loss: 2.4472, Validation accuracy: 0.47

//... epoch: 66/100
Training loss: 0.2709, Train accuracy: 0.93
Validation loss: 2.4706, Validation accuracy: 0.47

//... epoch: 67/100
Training loss: 0.2676, Train accuracy: 0.93
Validation loss: 2.4908, Validation accuracy: 0.47

//... epoch: 68/100
Training loss: 0.2603, Train accuracy: 0.93
Validation loss: 2.5064, Validation accuracy: 0.47

//... epoch: 69/100
Training loss: 0.2564, Train accuracy: 0.93
Validation loss: 2.5284, Validation accuracy: 0.47

//... epoch: 70/100
Training loss: 0.2550, Train accuracy: 0.94
Validation loss: 2.5394, Validation accuracy: 0.47

//... epoch: 71/100

Training loss: 0.2499, Train accuracy: 0.94
Validation loss: 2.5794, Validation accuracy: 0.47

//... epoch: 72/100
Training loss: 0.2430, Train accuracy: 0.94
Validation loss: 2.5728, Validation accuracy: 0.47

//... epoch: 73/100
Training loss: 0.2372, Train accuracy: 0.94
Validation loss: 2.6065, Validation accuracy: 0.47

//... epoch: 74/100
Training loss: 0.2367, Train accuracy: 0.94
Validation loss: 2.6145, Validation accuracy: 0.47

//... epoch: 75/100
Training loss: 0.2330, Train accuracy: 0.95
Validation loss: 2.6485, Validation accuracy: 0.47

//... epoch: 76/100
Training loss: 0.2280, Train accuracy: 0.95
Validation loss: 2.6715, Validation accuracy: 0.47

//... epoch: 77/100
Training loss: 0.2212, Train accuracy: 0.95
Validation loss: 2.6883, Validation accuracy: 0.47

//... epoch: 78/100
Training loss: 0.2163, Train accuracy: 0.95
Validation loss: 2.7061, Validation accuracy: 0.47

//... epoch: 79/100
Training loss: 0.2188, Train accuracy: 0.95
Validation loss: 2.7145, Validation accuracy: 0.47

//... epoch: 80/100
Training loss: 0.2111, Train accuracy: 0.95
Validation loss: 2.7535, Validation accuracy: 0.47

//... epoch: 81/100
Training loss: 0.2154, Train accuracy: 0.96
Validation loss: 2.7536, Validation accuracy: 0.47

//... epoch: 82/100
Training loss: 0.2042, Train accuracy: 0.96
Validation loss: 2.7587, Validation accuracy: 0.47

//... epoch: 83/100

Training loss: 0.1977, Train accuracy: 0.96
Validation loss: 2.7704, Validation accuracy: 0.47

//... epoch: 84/100
Training loss: 0.2032, Train accuracy: 0.96
Validation loss: 2.7826, Validation accuracy: 0.47

//... epoch: 85/100
Training loss: 0.1957, Train accuracy: 0.96
Validation loss: 2.7943, Validation accuracy: 0.47

//... epoch: 86/100
Training loss: 0.1932, Train accuracy: 0.96
Validation loss: 2.8058, Validation accuracy: 0.47

//... epoch: 87/100
Training loss: 0.1894, Train accuracy: 0.96
Validation loss: 2.8236, Validation accuracy: 0.47

//... epoch: 88/100
Training loss: 0.1879, Train accuracy: 0.96
Validation loss: 2.8286, Validation accuracy: 0.47

//... epoch: 89/100
Training loss: 0.1896, Train accuracy: 0.96
Validation loss: 2.8534, Validation accuracy: 0.46

//... epoch: 90/100
Training loss: 0.1869, Train accuracy: 0.96
Validation loss: 2.8501, Validation accuracy: 0.47

//... epoch: 91/100
Training loss: 0.1834, Train accuracy: 0.96
Validation loss: 2.8692, Validation accuracy: 0.47

//... epoch: 92/100
Training loss: 0.1785, Train accuracy: 0.96
Validation loss: 2.8754, Validation accuracy: 0.47

//... epoch: 93/100
Training loss: 0.1813, Train accuracy: 0.97
Validation loss: 2.8891, Validation accuracy: 0.47

//... epoch: 94/100
Training loss: 0.1812, Train accuracy: 0.96
Validation loss: 2.8988, Validation accuracy: 0.47

//... epoch: 95/100

Training loss: 0.1742, Train accuracy: 0.96
Validation loss: 2.9122, Validation accuracy: 0.47

//... epoch: 96/100
Training loss: 0.1719, Train accuracy: 0.97
Validation loss: 2.9151, Validation accuracy: 0.47

//... epoch: 97/100
Training loss: 0.1682, Train accuracy: 0.97
Validation loss: 2.9314, Validation accuracy: 0.47

//... epoch: 98/100
Training loss: 0.1720, Train accuracy: 0.97
Validation loss: 2.9471, Validation accuracy: 0.47

//... epoch: 99/100
Training loss: 0.1682, Train accuracy: 0.97
Validation loss: 2.9573, Validation accuracy: 0.46

//... epoch: 100/100
Training loss: 0.1633, Train accuracy: 0.97
Validation loss: 2.9596, Validation accuracy: 0.46


//...training the network took 896.53 sec...//

	precision	recall	f1-score	support
airplane	0.65	0.66	0.66	500
truck	0.50	0.52	0.51	500
bird	0.36	0.37	0.36	500
car	0.59	0.61	0.60	500
cat	0.28	0.28	0.28	500
deer	0.39	0.35	0.37	500
dog	0.28	0.28	0.28	500
horse	0.48	0.49	0.49	500
monkey	0.35	0.36	0.36	500
ship	0.63	0.59	0.61	500
accuracy			0.45	5000
macro avg	0.45	0.45	0.45	5000
weighted avg	0.45	0.45	0.45	5000

//...Confusion Matrix...//

```
[[331 33 20 23 5 18 10 7 3 50]
 [ 26 261 16 97 16 7 5 10 7 55]
 [ 28 5 185 21 59 57 45 23 68 9]
 [ 26 105 8 303 9 8 6 5 7 23]
 [ 4 20 48 8 139 59 85 36 83 18]
 [ 15 8 67 11 64 173 55 52 46 9]
```

```
[ 5  6 57  7 86 44 142 73 75  5]
[ 4 12 38  4 34 35 85 246 41  1]
[ 2  4 67 11 66 39 76 51 181  3]
[65 67  9 28 14  7  6  6  3 295]]
```



```
[87]: def imshow(img,title_grid=None):
    img = img / 2 + 0.5      # unnormalize
    npimg = img.numpy()
    label_name = {'1': 'airplane', '2': 'bird', '3': 'car', '4': 'cat', '5':
    → 'deer', '6': 'dog', '7': 'horse', '8': 'monkey', '9': 'ship', '10': 'truck'}
    tc=[]
    if title_grid is not None:
        for x in title_grid:
            a=label_name[x]
            tc.append(a)
        plt.title(tc)
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

classes=train_dataset.classes
testDataLoader_disp = DataLoader(test_dataset, batch_size=4,shuffle=True)
images_test,lab_test=next(iter(testDataLoader_disp))
title_lab=((classes[lab_test[j]]) for j in range(4))
```



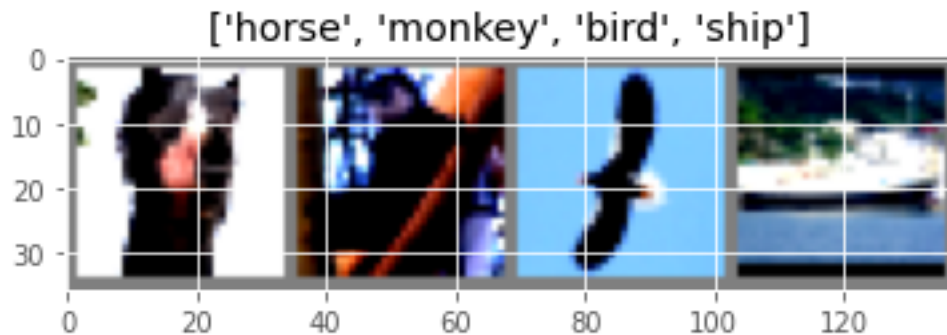
```

title_lab1=(list(title_lab))

imshow(torchvision.utils.make_grid(images_test),title_lab1)

```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



```

[115]: classes=train_dataset.classes
testDataLoader_disp = DataLoader(test_dataset, batch_size=1,shuffle=True)
model=LeNet5()
model.eval()
with torch.no_grad():
    image, label= next(iter(testDataLoader_disp))
    gtLabel = classes[label.numpy()[0]]
    pred = model(image)
    idx = pred.argmax(axis=1).cpu().numpy()[0]
    predLabel = classes[idx]
    if gtLabel != predLabel:
        print("ground truth label: {}, predicted label: {}".format(gtLabel,
→predLabel))
        imshow(torchvision.utils.make_grid(image),title_grid=None)

```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

ground truth label: 3, predicted label: 9

