

# Untitled

November 16, 2021

```
[510]: import numpy as np
import torch
import torch.nn as nn
from torchvision import models
import os, glob
from torch.utils.data import Dataset, DataLoader, random_split
from torchvision import transforms
from PIL import Image, ImageOps
import matplotlib.pyplot as plt
import torch.nn.functional as F
from torch.optim import Adam
import time
from datetime import datetime
```

```
[515]: def get_upsampling_weight(in_channels, out_channels, kernel_size):
    """Make a 2D bilinear kernel suitable for upsampling"""
    factor = (kernel_size + 1) // 2
    if kernel_size % 2 == 1:
        center = factor - 1
    else:
        center = factor - 0.5
    og = np.ogrid[:kernel_size, :kernel_size]
    filt = (1 - abs(og[0] - center) / factor) * \
           (1 - abs(og[1] - center) / factor)
    weight = np.zeros((in_channels, out_channels, kernel_size, kernel_size),
                      dtype=np.float64)
    weight[range(in_channels), range(out_channels), :, :] = filt
    return torch.from_numpy(weight).float()
```

```
[514]: class fcn32_resNet(nn.Module):
    def __init__(self):
        super().__init__()

        #initialize the pretrained model
        resNet18=models.resnet18(pretrained=True)

        #initialize the first 8 layers (0-7) of the pretrained model
        self.feature=nn.Sequential(*(list(resNet18.children())[0:8]))
```

```

#initialize the averagePooling layer
self.pool5=nn.AvgPool2d(kernel_size=7, stride=1, padding=0)

#initialize the FCN_1 layer
self.fcn6=nn.Conv2d(in_channels=512, out_channels=34, kernel_size=1)
self.relu6=nn.ReLU()

#intitalize the transpose conv layer
self.transConv7=nn.ConvTranspose2d(in_channels=34, out_channels=34,
↪kernel_size=64, stride=32)

self._initialize_weights()

def _initialize_weights(self):
    for m in self.modules():
        if isinstance(m, nn.ConvTranspose2d):
            assert m.kernel_size[0] == m.kernel_size[1]
            initial_weight = get_upsampling_weight(
                m.in_channels, m.out_channels, m.kernel_size[0])
            m.weight.data.copy_(initial_weight)

def forward(self,x):
    #print(x.size()[2])
    #passing the input through the layers of backbone network, ResNet18
    h=self.feature(x)
    #print("after",x.size()[2])
    #passing the previous output through the AvgPooling layer5
    h=self.pool5(h)

    #passing the previous output through the FCN layer
    h=self.fcn6(h)
    #h=self.relu6(h)

    #passing the previous output through the transpose conv layer
    h=self.transConv7(h)
    #print("h",h.shape)
    #print("h.size2",h.size()[2])
    #print("h.size3",h.size()[3])
    x_crop=(h.size()[2]-(x.size()[2]-200))//2
    y_crop=(h.size()[3]-(x.size()[3]-200))//2
    #print("x_crop",x_crop)
    #print("y_crop",y_crop)
    output = h[:, :, x_crop:x_crop+ x.size()[2]-200, y_crop:y_crop + x.
↪size()[3]-200].contiguous()

    return output

```

```

[516]: class fcn16_resNet(nn.Module):
    def __init__(self):
        super().__init__()

        #initialize the pretrained model
        resNet18=models.resnet18(pretrained=True)

        #initialize the first 7 layers (0-6) of the pretrained model
        self.feature1=nn.Sequential(*(list(resNet18.children())[0:7]))

        #initialize the 8th layer (7) of the pretrained model
        self.feature2=nn.Sequential(*(list(resNet18.children())[7:8]))

        #initialize the averagePooling layer
        self.pool5=nn.AvgPool2d(kernel_size=7, stride=1, padding=0)

        #initialize the fcn layer
        self.fcn6=nn.Conv2d(in_channels=512, out_channels=34, kernel_size=1)

        #initialize the tranpose conv layer on top of fcn layer
        self.upscore2 = nn.ConvTranspose2d(in_channels=34, out_channels=34,
↪kernel_size=4, stride=2)

        #initialize conv layer on top of pool4 layer
        self.convPool=nn.Conv2d(in_channels=256, out_channels=34, kernel_size=1)

        #initialize the transpose conv later on top of the new conv layer on
↪top of pool4 layer
        self.upscore16=nn.ConvTranspose2d(in_channels=34, out_channels=34,
↪kernel_size=32, stride=16)

        self._initialize_weights()

    def _initialize_weights(self):
        for m in self.modules():
            if isinstance(m, nn.ConvTranspose2d):
                assert m.kernel_size[0] == m.kernel_size[1]
                initial_weight = get_upsampling_weight(
                    m.in_channels, m.out_channels, m.kernel_size[0])
                m.weight.data.copy_(initial_weight)

    def forward(self,x):
        h=x

        #passing the input till the pool4 layers of backbone network, ResNet18
        h=self.feature1(h)

```

```

        #passing the previous output through the layer5
        h1=self.feature2(h)

        #passing the h1 output through the FCN layer
        h1=self.fcn6(h1)

        #passing the h1 output through the transpose conv layer to calculate ↪
        ↪upscore2
        h1=self.upscore2(h1)

        #passing the output of pool4 layer to another conv layer to predict new ↪
        ↪scores
        h2=self.convPool(h)

        #print("h2",h2.shape)
        pool4score = h2

        #print("poolscore",pool4score.shape)
        h1=h1[:, :, 1:1+h2.size()[2], 1:1+h2.size()[3]]
        #print("h1",h1.shape)
        h=h1+pool4score

        h=self.upscore16(h)
        x_crop=(h.size()[2]-(x.size()[2]-200))//2
        y_crop=(h.size()[3]-(x.size()[3]-200))//2
        output = h[:, :, x_crop:x_crop+ x.size()[2]-200, y_crop:y_crop + x.
        ↪size()[3]-200].contiguous()
        return output

```

```

[621]: color_dict={0:( 0,  0,  0),
                  1:( 0,  0,  0),
                  2:( 0,  0,  0),
                  3:( 0,  0,  0),
                  4:( 0,  0,  0),
                  5:(111,74,0),
                  6:(81,0,81),
                  7:(128,64,128),
                  8:(244,35,232),
                  9:(250,170,160),
                  10:(230,150,140),
                  11:(70,70,70),
                  12:(102,102,156),
                  13:(190,153,153),
                  14:(180,165,180),
                  15:(150,100,100),
                  16:(150,120,90),

```

```

17: (153, 153, 153),
18: (153, 153, 153),
19: (250, 170, 30),
20: (220, 220, 0),
21: (107, 142, 35),
22: (152, 251, 152),
23: (70, 130, 180),
24: (220, 20, 60),
25: (255, 0, 0),
26: (0, 0, 142),
27: (0, 0, 70),
28: (0, 60, 100),
29: (0, 0, 90),
30: (0, 0, 110),
31: (0, 80, 100),
32: (0, 0, 230),
33: (119, 11, 32)}

```

```
color_dict[33]
```

```
[621]: (119, 11, 32)
```

```
[517]: res=models.resnet18(pretrained=True)
```

```
[361]: feat=nn.Sequential(*(list(res.children())[0:8]))
#feat
```

```
[475]: model_fcn32=fcn32_resNet()
#model_fcn32
```

```
[456]: model_fcn16=fcn16_resNet()
#model_fcn16
```

```
[518]: a="C:\\Users\\hp631\\USC_Jupyter\\Fall'21\\Advanced computer_
↳vision\\HW5\\data\\datafile\\training\\image_2"

#a=os.listdir(r"C:\Users\hp631\USC_Jupyter\Fall'21\Advanced computer_
↳vision\HW5\data\datafile\training\image_2")
#b=os.listdir(r"C:\Users\hp631\USC_Jupyter\Fall'21\Advanced computer_
↳vision\HW5\data\datafile\training\semantic")
b="C:\\Users\\hp631\\USC_Jupyter\\Fall'21\\Advanced computer_
↳vision\\HW5\\data\\datafile\\training\\semantic"
print(b)
```

```
C:\Users\hp631\USC_Jupyter\Fall'21\Advanced computer
vision\HW5\data\datafile\training\semantic
```

```
[520]: class KITTI_custom(Dataset):
    def __init__(self,
        ↳image_dir,mask_dir,transform_img=None,transform_mask=None):
        self.transform_img=transform_img
        self.transform_mask=transform_mask
        self.image_dir=image_dir
        self.mask_dir=mask_dir
        self.img_files=os.listdir(image_dir)
        self.mask_files=os.listdir(mask_dir)
        self.datalist=[]
        for i in range(len(self.img_files)):
            imgPath=os.path.join(self.image_dir,self.img_files[i])
            maskPath=os.path.join(self.mask_dir,self.mask_files[i])
            image=Image.open(imgPath)
            mask=Image.open(maskPath)
            image=ImageOps.expand(image, border = 100, fill = 0)
            if self.transform_img is not None:
                x=self.transform_img(image)
            else:
                x=image
            if self.transform_mask is not None:
                y=self.transform_mask(mask)
            else:
                y=mask
            self.datalist.append((x,y))

    def __getitem__(self, index):
        return self.datalist[index]

    def __len__(self):
        return len(self.img_files)
```

```
[521]: mydataset=KITTI_custom(a,b,transforms.ToTensor(),transforms.Compose([transforms.
    ↳ToTensor(), lambda x: x*255]))
dataloader=DataLoader(mydataset,batch_size=1,shuffle=True)
x1,y1=next(iter(dataloader))
plt.imshow((x1.squeeze()).permute(1,2,0))
#plt.imshow(y1.squeeze())
print(x1.shape)
print(x1)
```

```
torch.Size([1, 3, 575, 1442])
tensor([[[[0., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 0.],
          ...,
          [0., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 0.],
```

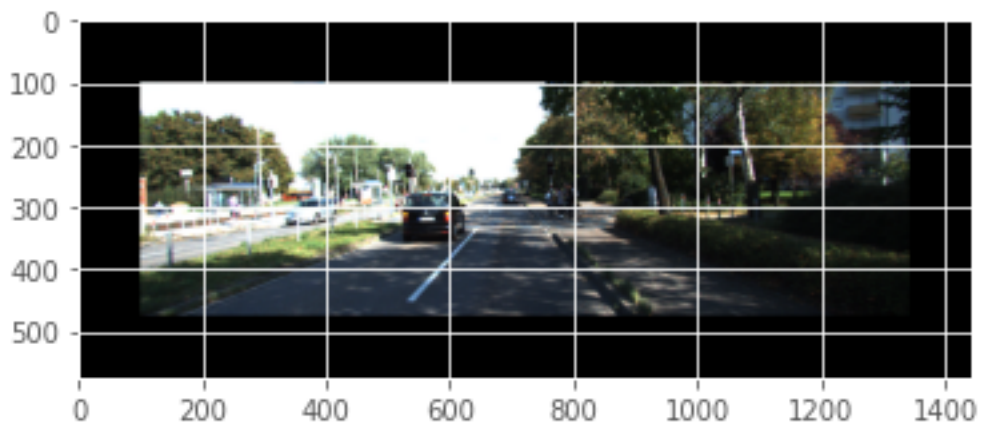
```

[0., 0., 0., ..., 0., 0., 0.],

[[0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 ...,
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.]],


[[0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 ...,
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.]])

```



```
[522]: plt.imshow(y1.squeeze())
```

```
[522]: <matplotlib.image.AxesImage at 0x2001680f670>
```



```
[523]: classes=np.unique(y1.detach().numpy())
       classes
```

```
[523]: array([ 4.,  7.,  8., 10., 11., 13., 17., 19., 20., 21., 22., 23., 24.,
        25., 26., 28., 32., 33.], dtype=float32)
```

```
[525]: #split dataset into train, validation, test
       train_size=int(0.7*len(mydataset))
       val_size=int(0.15*len(mydataset))
       test_size=len(mydataset)-(train_size+val_size)
       train_dataset,
       ↪val_dataset,test_dataset=random_split(mydataset,[train_size,val_size,test_size])
       print("Train size",len(train_dataset))
       print("Validation size",len(val_dataset))
       print("test size",len(test_dataset))
```

Train size 140

Validation size 30

test size 30

```
[533]: def
       ↪experiment(train_data,val_data,epochs,model_version,batchsize,learning_rate):

       # initializing the train, validation, and test data loaders with given
       ↪batchsize
       trainDataLoader = DataLoader(train_data, shuffle=False,batch_size=batchsize)
       valDataLoader = DataLoader(val_data, batch_size=batchsize)

       # steps per epoch for train and validation set
       trainSteps = len(trainDataLoader.dataset) // batchsize
       valSteps = len(valDataLoader.dataset) // batchsize

       #initialize the model
       print("//...Initializing {} model...//".format(model_version))
       #model= FCN32/16()
       model=model_version()

       #defining optimizer and cross-entropy loss function
       optimizer=Adam(model.parameters(),lr=learning_rate)
       criterion=nn.CrossEntropyLoss()

       #Learning rate will reduce by factor of 25% after every step size of
       ↪15epochs
       scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=15,
       ↪gamma=0.25)
```



```

#empty dict to store losses and accuracy score
myDict={"train_loss": [], "val_loss": []}

#initialize training the model
print("//...Training initiated...//")
start=time.time()

for epo in range(0,epochs):
    model.train()
    trainLoss=0
    valLoss=0
    for (x,y) in trainDataLoader:
        optimizer.zero_grad()
        ytrain_pred=model(x)
        target=y.long().squeeze(1)
        loss=criterion(ytrain_pred,target)
        loss.backward()
        optimizer.step()
        trainLoss+=loss
    with torch.no_grad():
        model.eval()
        for (x,y) in valDataLoader:
            yval_pred=model(x)
            target_val=y.long().squeeze(1)
            loss=criterion(yval_pred,target_val)
            valLoss+=loss
    #scheduler.step()

    #calculating loss and accuracy for each epoch for both train and val
    ↪sets
    trainLoss_avg=trainLoss/trainSteps
    valLoss_avg=valLoss/valSteps

    myDict["train_loss"].append(trainLoss_avg.cpu().detach().numpy())
    myDict["val_loss"].append(valLoss_avg.cpu().detach().numpy())

    print("//... epoch: {}/{}".format(epo + 1, epochs))
    print("Training loss: {:.4f}".format(trainLoss_avg))
    print("Validation loss: {:.4f} \n".format(valLoss_avg))

end=time.time()
print("//...training the network took {:.2f} sec...//".format(end-start))

#plotting loss and acc
plt.style.use("ggplot")
plt.figure()

```

```

plt.plot(myDict["train_loss"], label="train_loss")
plt.plot(myDict["val_loss"], label="val_loss")
plt.title("Training Loss : {} model".format(type(model).__name__))
plt.xlabel("# of Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()

return model

```

```

[534]: model_32=experiment(train_dataset,val_dataset,20,fcn32_resNet,1,0.001)
PATH="E:\\usc\\Fall'21\\CSCI677 Adv Computer Vision\\HW5\\models\\model32\\{}.
→pth".format(datetime.now().strftime("%d_%m_%H_%M"))
torch.save(model_32.state_dict(), PATH)

```

```

//...Initializing <class '__main__.fcn32_resNet'> model...//
//...Training initiated...//
//... epoch: 1/20
Training loss: 1.3115
Validation loss: 1.2198

//... epoch: 2/20
Training loss: 1.0315
Validation loss: 1.1509

//... epoch: 3/20
Training loss: 0.9084
Validation loss: 1.0056

//... epoch: 4/20
Training loss: 0.8042
Validation loss: 1.0201

//... epoch: 5/20
Training loss: 0.7479
Validation loss: 0.8942

//... epoch: 6/20
Training loss: 0.7018
Validation loss: 1.0156

//... epoch: 7/20
Training loss: 0.6481
Validation loss: 0.7914

//... epoch: 8/20
Training loss: 0.6181
Validation loss: 0.7717

```

//... epoch: 9/20  
Training loss: 0.5757  
Validation loss: 0.8414

//... epoch: 10/20  
Training loss: 0.5883  
Validation loss: 0.9546

//... epoch: 11/20  
Training loss: 0.5483  
Validation loss: 0.8399

//... epoch: 12/20  
Training loss: 0.4992  
Validation loss: 0.7460

//... epoch: 13/20  
Training loss: 0.4779  
Validation loss: 0.7603

//... epoch: 14/20  
Training loss: 0.4746  
Validation loss: 0.7725

//... epoch: 15/20  
Training loss: 0.4720  
Validation loss: 0.7558

//... epoch: 16/20  
Training loss: 0.4426  
Validation loss: 0.7623

//... epoch: 17/20  
Training loss: 0.4319  
Validation loss: 0.7075

//... epoch: 18/20  
Training loss: 0.4209  
Validation loss: 0.7409

//... epoch: 19/20  
Training loss: 0.4137  
Validation loss: 0.7331

//... epoch: 20/20  
Training loss: 0.4045  
Validation loss: 0.8215

//...training the network took 8822.77 sec...//



```
[535]: model_16=experiment(train_dataset,val_dataset,20,fcn16_resNet,1,0.001)
PATH="E:\\usc\\Fall'21\\CSCI677 Adv Computer Vision\\HW5\\models\\model16\\{}.
.pth".format(datetime.now().strftime("%d_%m_%H_%M"))
torch.save(model_16.state_dict(), PATH)
```

//...Initializing <class '\_\_main\_\_.fcn16\_resNet'> model...//

//...Training initiated...//

//... epoch: 1/20

Training loss: 1.1172

Validation loss: 1.0747

//... epoch: 2/20

Training loss: 0.7545

Validation loss: 0.8471

//... epoch: 3/20

Training loss: 0.6016

Validation loss: 0.7293

//... epoch: 4/20

Training loss: 0.5235

Validation loss: 0.8314

//... epoch: 5/20

Training loss: 0.4876

Validation loss: 0.8042

//... epoch: 6/20

Training loss: 0.4331

Validation loss: 0.7126

//... epoch: 7/20

Training loss: 0.3862

Validation loss: 0.6207

//... epoch: 8/20

Training loss: 0.3594

Validation loss: 0.6130

//... epoch: 9/20

Training loss: 0.2802

Validation loss: 0.6541

//... epoch: 10/20

Training loss: 0.2443

Validation loss: 0.5291

//... epoch: 11/20

Training loss: 0.2257

Validation loss: 0.6113

//... epoch: 12/20

Training loss: 0.2502

Validation loss: 0.6484

//... epoch: 13/20

Training loss: 0.2388

Validation loss: 0.5536

//... epoch: 14/20

Training loss: 0.2207

Validation loss: 0.5699

//... epoch: 15/20

Training loss: 0.2085

Validation loss: 0.5940

//... epoch: 16/20

Training loss: 0.1950

Validation loss: 0.5723

//... epoch: 17/20

Training loss: 0.1684

Validation loss: 0.5364

//... epoch: 18/20

Training loss: 0.1482

Validation loss: 0.5846

//... epoch: 19/20

Training loss: 0.1377

Validation loss: 0.5444

//... epoch: 20/20

Training loss: 0.1308

Validation loss: 0.5736

//...training the network took 9862.73 sec...//



```
[641]: test_dataloader=DataLoader(test_dataset,batch_size=1,shuffle=True)
x1,y1=next(iter(test_dataloader))
#plt.imshow((x1.squeeze()).permute(1,2,0))
```

```

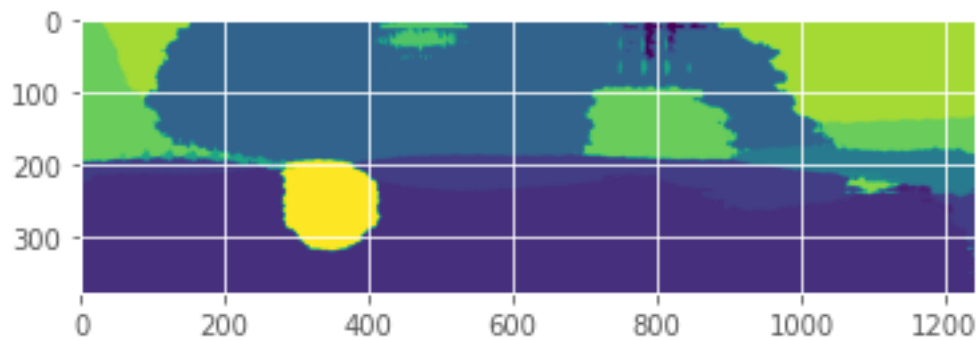
plt.imshow(y1.squeeze())
#print(x1.shape)
#print(x1)

y_pred_32=model_32.forward(x1)
#print(y_pred_32.shape)
#print(y1.shape)
y_pred_32_lab=np.argmax(((y_pred_32.permute(0,2,3,1)).detach().
    ↪numpy()),axis=3)[0,:,:]
#y_pred_32_soft_v2=np.argmax((y_pred_32.detach().numpy()),axis=1)[0,:,:]

#print(y_pred_32_lab.shape)
#print(y_pred_32_lab)
#print(y_pred_32)
#print(y1)
plt.imshow(y1.squeeze())
plt.imshow(y_pred_32_lab)


```

[641]: <matplotlib.image.AxesImage at 0x2001b3e28e0>




[642]: plt.imshow(y1.squeeze())

[642]: <matplotlib.image.AxesImage at 0x2001b449640>



```
[643]: plt.imshow((x1.squeeze()).permute(1,2,0))
```

```
[643]: <matplotlib.image.AxesImage at 0x2002de4a880>
```




```
[644]: y_pred_16=model_16.forward(x1)
        #print(y_pred_16.shape)
        #print(y1.shape)
        y_pred_16_lab=np.argmax(((y_pred_16.permute(0,2,3,1)).detach().
        ↪numpy()),axis=3)[0,:,:]
        #y_pred_32_soft_v2=np.argmax((y_pred_32.detach().numpy()),axis=1)[0,:,:]

        #print(y_pred_16_lab.shape)
        #print(y_pred_16_lab)
        #print(y_pred_16)
        #print(y1)
        #plt.imshow(y1.squeeze())
        plt.imshow(y_pred_16_lab)
```


```
[644]: <matplotlib.image.AxesImage at 0x2002e0c0130>
```





```
[645]: pred = (torch.argmax(y_pred_16, dim=1)).squeeze().numpy()
color_img=np.zeros((pred.shape[0],pred.shape[1],3))


for i in range(pred.shape[0]):
    for j in range(pred.shape[1]):
        rgb_val=list(color_dict[pred[i][j]])
        for k in range(3):
            color_img[i][j][k]=rgb_val[k]
color_img=color_img.astype(np.uint8)
plt.imshow(color_img)
plt.show()
```



```
[646]: pred_32 = (torch.argmax(y_pred_32, dim=1)).squeeze().numpy()
color_img_32=np.zeros((pred_32.shape[0],pred_32.shape[1],3))

for i in range(pred_32.shape[0]):
    for j in range(pred_32.shape[1]):
        rgb_val=list(color_dict[pred_32[i][j]])
        for k in range(3):
            color_img_32[i][j][k]=rgb_val[k]
```

```
color_img_32=color_img_32.astype(np.uint8)
plt.imshow(color_img_32)
plt.show()
```



```
[659]: #dummy code
ypred=(torch.argmax(model_16.forward(x1), dim=1)).squeeze().numpy()
ytrue=y1.squeeze().detach().numpy()
print(ypred.shape[0])
print(ytrue.shape)
```

```
376
(376, 1241)
```

```
[700]: #dummy code
a1=np.array([[0,1,0],[1,0,1]])
a2=np.array([[0,1,0],[0,1,1]])
intersection=np.sum(np.abs(a1*a2),axis=(0,1))
mask_sum = np.sum(np.abs(a1), axis=(0,1)) + np.sum(np.abs(a2), axis=(0,1))
union = mask_sum - intersection
iou_temp=intersection/union
iou_temp
```

```
[700]: 0.5
```

```
[663]: def oneHot(Ypred,Ytrue,label_pos):
    ypred_hot=np.zeros((Ypred.shape[0],Ypred.shape[1]))
    ytrue_hot=np.zeros((Ytrue.shape[0],Ytrue.shape[1]))
    for i in range(Ypred.shape[0]):
        for j in range(Ypred.shape[1]):
            if Ypred[i][j]==label_pos:
                ypred_hot[i][j]=1
            else:
                ypred_hot[i][j]=0
            if Ytrue[i][j]==label_pos:
```

```

        ytrue_hot[i][j]=1
    else:
        ytrue_hot[i][j]=0
    return ypred_hot,ytrue_hot

```

```

[668]: testDataLoader = DataLoader(test_dataset, shuffle=False, batch_size=1)
smooth=0.001
iou=[]
for cls in range(34):
    intersection_class=0
    union_class=0
    with torch.no_grad():
        model_16.eval()
        for (x,y) in testDataLoader:
            ypred=(torch.argmax(model_16.forward(x), dim=1)).squeeze().numpy()
            ytrue=y.squeeze().detach().numpy()
            ypred_hot,ytrue_hot=oneHot(ypred,ytrue,cls)
            intersection=np.sum(np.abs(ypred_hot*ytrue_hot),axis=(0,1))
            mask_sum = np.sum(np.abs(ytrue_hot), axis=(0,1)) + np.sum(np.
↪abs(ypred_hot), axis=(0,1))
            union = mask_sum - intersection
            intersection_class=intersection_class+intersection
            union_class=union_class+union
        score=(intersection_class+smooth)/(union_class+smooth)
        iou.append(score)

```

```

[719]: iou_scores=iou
for i,sco in enumerate(iou_scores):
    if sco==1:
        iou_scores[i]='N/A'
iou_scores

```

```

[719]: [2.2222172839615914e-06,
        'N/A',
        'N/A',
        'N/A',
        0.231036164739202,
        0.050636755358297175,
        0.09985620065583961,
        0.8368964067061173,
        0.37932320877778863,
        0.047371263515181505,
        0.706572502004664,
        0.6594890555327194,
        0.34769565951593834,
        0.13958343161002015,


```

```
0.4351507003795953,  
0.0878950441202621,  
3.3394668095753643e-09,  
0.32514361258903973,  
9.615292160652301e-06,  
0.4087523506846543,  
0.2187025981908527,  
0.7931750532250866,  
0.7833640451679014,  
0.9276868336407721,  
0.08001980541315529,  
8.77962354730154e-07,  
0.7512057936058436,  
0.20530787972981865,  
2.5562371534704206e-08,  
2.6130121209793257e-07,  
6.353236116114285e-07,  
1.1650936543057608e-07,  
2.7570987711335064e-07,  
0.07022727707584585]
```

```
[712]: scores_16=[]  
for sco in iou_scores:  
    if sco!='N/A':  
        scores_16.append(sco)  
miou_16=np.mean(scores_16)  
print("mIoU for FCN16 netowrk = {}".format(miou_16))
```

mIoU for FCN16 netowrk = 0.2769388927566549

```
[718]: x_scale=[]  
for j in range(34):  
    x_scale.append("class"+str(j))  
  
fig, ax = plt.subplots(figsize=(16, 9))  
ax.barh(x_scale, iou, align='center')  
ax.invert_yaxis()  
plt.show()
```



```
[684]: testDataLoader = DataLoader(test_dataset, shuffle=False, batch_size=1)
smooth=0.001
iou_32=[]
for cls in range(34):
    intersection_class=0
    union_class=0
    with torch.no_grad():
        model_32.eval()
        for (x,y) in testDataLoader:
            ypred=(torch.argmax(model_32.forward(x), dim=1)).squeeze().numpy()
            ytrue=y.squeeze().detach().numpy()
            ypred_hot,ytrue_hot=oneHot(ypred,ytrue,cls)
            intersection=np.sum(np.abs(ypred_hot*ytrue_hot),axis=(0,1))
            mask_sum = np.sum(np.abs(ytrue_hot), axis=(0,1)) + np.sum(np.
↪abs(ypred_hot), axis=(0,1))
            union = mask_sum - intersection
            intersection_class=intersection_class+intersection
            union_class=union_class+union
        score=(intersection_class+smooth)/(union_class+smooth)
        iou_32.append(score)
```


```
[706]: iou_scores_32=iou_32
for i,sco in enumerate(iou_scores_32):
    if sco==1:
        iou_scores_32[i]='N/A'
iou_scores_32
```

```
[706]: [2.4096327478728968e-06,
        'N/A',
        'N/A',
        'N/A',
        0.12957794552844612,
        1.5119441318499953e-07,
        0.08132142504381275,
        0.8081879788867206,
        0.2919028296615624,
        2.31481476123114e-08,
        0.31895706817178393,
        0.5546574405534996,
        0.45665644148383766,
        0.1055692507591559,
        0.20401999038000185,
        4.659830074636498e-07,
        3.340281106901418e-09,
        0.027847473597112257,
        9.615292160652301e-06,
        0.00279012790127491,
        2.4074340987184243e-08,
        0.7076955413350989,
        0.5958109202161124,
        0.8116267771882951,
        7.942180292099096e-08,
        1.6286618425702891e-06,
        0.5343571795753566,
        0.005967691589713594,
        2.5562371534704206e-08,
        4.208752437393755e-07,
        7.037292725339391e-07,
        0.010255546312181027,
        2.7570987711335064e-07,
        4.0584399113474384e-07]
```

```
[710]: scores=[]
        for sco in iou_scores_32:
            if sco!='N/A':
                scores.append(sco)
        miou_32=np.mean(scores)
        print("mIoU for FCN32 netowrk = {}".format(miou_32))
```

mIoU for FCN32 netowrk = 0.1821683180855957

```
[686]: fig, ax = plt.subplots(figsize =(16, 9))
        ax.barh(x_scale, iou_scores_32, align='center')
        ax.invert_yaxis()
        plt.show()
```



```
[458]: #dummy code
loss = nn.CrossEntropyLoss()
inp=y_pred_32
target=y1.long().squeeze(1)
target_sub=target-1
print("input",inp.shape)
print("target",target.shape)
print("target_sub",target_sub.shape)
print("input",inp)
print("target",target)
print("target_sub",target_sub)

output = loss(inp,target_sub)
print("loss",output)

input torch.Size([1, 34, 374, 1238])
target torch.Size([1, 374, 1238])
target_sub torch.Size([1, 374, 1238])
input tensor([[[[ 0.0014,  0.0014,  0.0014, ...,  0.0014,  0.0014,  0.0014],
                  [ 0.0014,  0.0014,  0.0014, ...,  0.0014,  0.0014,  0.0014],
                  [ 0.0014,  0.0014,  0.0014, ...,  0.0014,  0.0014,  0.0014],
                  ...,
                  [ 0.0014,  0.0014,  0.0014, ...,  0.0014,  0.0014,  0.0014],
                  [ 0.0014,  0.0014,  0.0014, ...,  0.0014,  0.0014,  0.0014],
                  [ 0.0014,  0.0014,  0.0014, ...,  0.0014,  0.0014,  0.0014]]],
               [[ 0.0014,  0.0014,  0.0014, ...,  0.0014,  0.0014,  0.0014],
```

```

[ 0.0014, 0.0014, 0.0014, ..., 0.0014, 0.0014, 0.0014],
[ 0.0014, 0.0014, 0.0014, ..., 0.0014, 0.0014, 0.0014],
...,
[ 0.0014, 0.0014, 0.0014, ..., 0.0014, 0.0014, 0.0014],
[ 0.0014, 0.0014, 0.0014, ..., 0.0014, 0.0014, 0.0014],
[ 0.0014, 0.0014, 0.0014, ..., 0.0014, 0.0014, 0.0014]],

[[-0.0012, -0.0012, -0.0012, ..., -0.0012, -0.0012, -0.0012],
[-0.0012, -0.0012, -0.0012, ..., -0.0012, -0.0012, -0.0012],
[-0.0012, -0.0012, -0.0012, ..., -0.0012, -0.0012, -0.0012],
...,
[-0.0012, -0.0012, -0.0012, ..., -0.0012, -0.0012, -0.0012],
[-0.0012, -0.0012, -0.0012, ..., -0.0012, -0.0012, -0.0012],
[-0.0012, -0.0012, -0.0012, ..., -0.0012, -0.0012, -0.0012]],

...,

[[ 0.0015, 0.0015, 0.0015, ..., 0.0015, 0.0015, 0.0015],
[ 0.0015, 0.0015, 0.0015, ..., 0.0015, 0.0015, 0.0015],
[ 0.0015, 0.0015, 0.0015, ..., 0.0015, 0.0015, 0.0015],
...,
[ 0.0015, 0.0015, 0.0015, ..., 0.0015, 0.0015, 0.0015],
[ 0.0015, 0.0015, 0.0015, ..., 0.0015, 0.0015, 0.0015],
[ 0.0015, 0.0015, 0.0015, ..., 0.0015, 0.0015, 0.0015]],

[[-0.0013, -0.0013, -0.0013, ..., -0.0013, -0.0013, -0.0013],
[-0.0013, -0.0013, -0.0013, ..., -0.0013, -0.0013, -0.0013],
[-0.0013, -0.0013, -0.0013, ..., -0.0013, -0.0013, -0.0013],
...,
[-0.0013, -0.0013, -0.0013, ..., -0.0013, -0.0013, -0.0013],
[-0.0013, -0.0013, -0.0013, ..., -0.0013, -0.0013, -0.0013],
[-0.0013, -0.0013, -0.0013, ..., -0.0013, -0.0013, -0.0013]],

[[ 0.0008, 0.0008, 0.0008, ..., 0.0008, 0.0008, 0.0008],
[ 0.0008, 0.0008, 0.0008, ..., 0.0008, 0.0008, 0.0008],
[ 0.0008, 0.0008, 0.0008, ..., 0.0008, 0.0008, 0.0008],
...,
[ 0.0008, 0.0008, 0.0008, ..., 0.0008, 0.0008, 0.0008],
[ 0.0008, 0.0008, 0.0008, ..., 0.0008, 0.0008, 0.0008],
[ 0.0008, 0.0008, 0.0008, ..., 0.0008, 0.0008, 0.0008]]],
grad_fn=<CopyBackwards>)
target tensor([[[[11, 11, 11, ..., 21, 21, 21],
[11, 11, 11, ..., 21, 21, 21],
[11, 11, 11, ..., 21, 21, 21],
...,
[ 7, 7, 7, ..., 7, 7, 7],
[ 7, 7, 7, ..., 7, 7, 7],
[ 7, 7, 7, ..., 7, 7, 7]]]])

```



```
target_sub tensor([[[10, 10, 10, ..., 20, 20, 20],
                    [10, 10, 10, ..., 20, 20, 20],
                    [10, 10, 10, ..., 20, 20, 20],
                    ...,
                    [ 6,  6,  6, ...,  6,  6,  6],
                    [ 6,  6,  6, ...,  6,  6,  6],
                    [ 6,  6,  6, ...,  6,  6,  6]])
loss tensor(3.5257, grad_fn=<NllLoss2DBackward>)
```

[ ]:

[ ]: