

```
!pip install pyyaml==5.1
```

```
import torch
```

```
TORCH_VERSION = ".".join(torch.__version__.split(".")[ :2])
```

```
CUDA_VERSION = torch.__version__.split("+")[-1]
```

```
print("torch: ", TORCH_VERSION, "; cuda: ", CUDA_VERSION)
```

```
# Install detectron2 that matches the above pytorch version
```

```
# See https://detectron2.readthedocs.io/tutorials/install.html for instructions
```

```
!pip install detectron2 -f https://dl.fbaipublicfiles.com/detectron2/wheels/\$CUDA\_VERSION/tor
```

```
# If there is not yet a detectron2 release that matches the given torch + CUDA version, you r
```

```

Collecting pyyaml==5.1
  Downloading PyYAML-5.1.tar.gz (274 kB)
    |████████████████████████████████████████| 274 kB 5.4 MB/s
Building wheels for collected packages: pyyaml
  Building wheel for pyyaml (setup.py) ... done
  Created wheel for pyyaml: filename=PyYAML-5.1-cp37-cp37m-linux_x86_64.whl size=44092
  Stored in directory: /root/.cache/pip/wheels/77/f5/10/d00a2bd30928b972790053b5de0c703
Successfully built pyyaml
Installing collected packages: pyyaml
  Attempting uninstall: pyyaml
    Found existing installation: PyYAML 3.13
    Uninstalling PyYAML-3.13:
      Successfully uninstalled PyYAML-3.13
Successfully installed pyyaml-5.1
torch: 1.10 ; cuda: cu111
Looking in links: https://dl.fbaipublicfiles.com/detectron2/wheels/cu111/torch1.10/index.html
Collecting detectron2
  Downloading https://dl.fbaipublicfiles.com/detectron2/wheels/cu111/torch1.10/detectron2-0.1.1-cp37-cp37m-linux\_x86\_64.whl
    |████████████████████████████████████████| 7.0 MB 759 kB/s
Collecting omegaconf>=2.1
  Downloading omegaconf-2.1.1-py3-none-any.whl (74 kB)
    |████████████████████████████████████████| 74 kB 2.3 MB/s
Requirement already satisfied: tensorboard in /usr/local/lib/python3.7/dist-packages (from detectron2)
Requirement already satisfied: pydot in /usr/local/lib/python3.7/dist-packages (from detectron2)
Collecting black==21.4b2
  Downloading black-21.4b2-py3-none-any.whl (130 kB)
    |████████████████████████████████████████| 130 kB 11.1 MB/s
Requirement already satisfied: Pillow>=7.1 in /usr/local/lib/python3.7/dist-packages (from detectron2)
Requirement already satisfied: pycocotools>=2.0.2 in /usr/local/lib/python3.7/dist-packages (from detectron2)
Requirement already satisfied: cloudpickle in /usr/local/lib/python3.7/dist-packages (from detectron2)
Requirement already satisfied: termcolor>=1.1 in /usr/local/lib/python3.7/dist-packages (from detectron2)
Requirement already satisfied: tabulate in /usr/local/lib/python3.7/dist-packages (from detectron2)
Collecting iopath<0.1.10,>=0.1.7
  Downloading iopath-0.1.9-py3-none-any.whl (27 kB)
Collecting yacs>=0.1.8
  Downloading yacs-0.1.8-py3-none-any.whl (14 kB)
Requirement already satisfied: tqdm>4.29.0 in /usr/local/lib/python3.7/dist-packages (from detectron2)
Requirement already satisfied: future in /usr/local/lib/python3.7/dist-packages (from detectron2)
Collecting fvc<0.1.6,>=0.1.5
  Downloading fvc-0.1.5.post20211023.tar.gz (49 kB)
    |████████████████████████████████████████| 49 kB 5.6 MB/s
Collecting hydra-core>=1.1
  Downloading hydra_core-1.1.1-py3-none-any.whl (145 kB)
    |████████████████████████████████████████| 145 kB 38.9 MB/s
Requirement already satisfied: matplotlib in /usr/local/lib/python3.7/dist-packages (from detectron2)
Requirement already satisfied: appdirs in /usr/local/lib/python3.7/dist-packages (from detectron2)
Collecting regex>=2020.1.8
  Downloading regex-2021.11.10-cp37-cp37m-manylinux_2_17_x86_64.manylinux2014_x86_64.whl
    |████████████████████████████████████████| 749 kB 47.2 MB/s
Collecting mpyc-extensions>=0.4.3
  Downloading mpyc_extensions-0.4.3-py2.py3-none-any.whl (4.5 kB)
Collecting typed-ast>=1.4.2
  Downloading typed_ast-1.5.0-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.manylin
    |████████████████████████████████████████| 843 kB 26.3 MB/s
Requirement already satisfied: toml>=0.10.1 in /usr/local/lib/python3.7/dist-packages (from detectron2)
Requirement already satisfied: typing-extensions>=3.7.4 in /usr/local/lib/python3.7/dist-packages (from detectron2)
Collecting pathspec<1,>=0.8.1

```

```

Downloading pathspec-0.9.0-py2.py3-none-any.whl (31 kB)
Requirement already satisfied: click>=7.1.2 in /usr/local/lib/python3.7/dist-packages (
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from fv
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.7/dist-packages (f
Requirement already satisfied: importlib-resources in /usr/local/lib/python3.7/dist-pac
Collecting antlr4-python3-runtime==4.8

```

```

Downloading antlr4-python3-runtime-4.8.tar.gz (112 kB)

```

```

| 112 kB 51.3 MB/s

```

```

Collecting portalocker

```

```

Downloading portalocker-2.3.2-py2.py3-none-any.whl (15 kB)

```

```

Requirement already satisfied: cython>=0.27.3 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: setuptools>=18.0 in /usr/local/lib/python3.7/dist-packag
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.7/dist-pa
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.7/dist-packa
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.7/dist-packages (
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/l
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-packages (from
Requirement already satisfied: zipp>=3.1.0 in /usr/local/lib/python3.7/dist-packages (f
Requirement already satisfied: tensorboard-data-server<0.7.0,>=0.6.0 in /usr/local/lib/
Requirement already satisfied: protobuf>=3.6.0 in /usr/local/lib/python3.7/dist-package
Requirement already satisfied: absl-py>=0.4 in /usr/local/lib/python3.7/dist-packages (

```

```

import detectron2

```

```

from detectron2.utils.logger import setup_logger

```

```

setup_logger()

```

```

# import some common libraries

```

```

import numpy as np

```

```

import os, json, cv2, random

```

```

from google.colab.patches import cv2_imshow

```

```

# import some common detectron2 utilities

```

```

from detectron2 import model_zoo

```

```

from detectron2.engine import DefaultPredictor

```

```

from detectron2.config import get_cfg

```

```

from detectron2.utils.visualizer import Visualizer

```

```

from detectron2.data import MetadataCatalog, DatasetCatalog

```

```

from detectron2.evaluation import PascalVOCDetectionEvaluator

```

```

!wget http://host.robots.ox.ac.uk/pascal/VOC/voc2007/VOCtrainval_06-Nov-2007.tar

```

```

!tar -xvf VOCtrainval_06-Nov-2007.tar

```

```

VOCdevkit/VOC2007/SegmentationObject/008722.png
VOCdevkit/VOC2007/SegmentationObject/008747.png
VOCdevkit/VOC2007/SegmentationObject/008764.png
VOCdevkit/VOC2007/SegmentationObject/008801.png
VOCdevkit/VOC2007/SegmentationObject/008815.png
VOCdevkit/VOC2007/SegmentationObject/008927.png
VOCdevkit/VOC2007/SegmentationObject/008932.png
VOCdevkit/VOC2007/SegmentationObject/008944.png
VOCdevkit/VOC2007/SegmentationObject/008948.png
VOCdevkit/VOC2007/SegmentationObject/008973.png
VOCdevkit/VOC2007/SegmentationObject/008980.png
VOCdevkit/VOC2007/SegmentationObject/009015.png
VOCdevkit/VOC2007/SegmentationObject/009068.png
VOCdevkit/VOC2007/SegmentationObject/009209.png
VOCdevkit/VOC2007/SegmentationObject/009221.png

```

```

VOCdevkit/VOC2007/SegmentationObject/009221.png
VOCdevkit/VOC2007/SegmentationObject/009245.png
VOCdevkit/VOC2007/SegmentationObject/009251.png
VOCdevkit/VOC2007/SegmentationObject/009252.png
VOCdevkit/VOC2007/SegmentationObject/009295.png
VOCdevkit/VOC2007/SegmentationObject/009323.png
VOCdevkit/VOC2007/SegmentationObject/009327.png
VOCdevkit/VOC2007/SegmentationObject/009331.png
VOCdevkit/VOC2007/SegmentationObject/009348.png
VOCdevkit/VOC2007/SegmentationObject/009392.png
VOCdevkit/VOC2007/SegmentationObject/009413.png
VOCdevkit/VOC2007/SegmentationObject/009419.png
VOCdevkit/VOC2007/SegmentationObject/009422.png
VOCdevkit/VOC2007/SegmentationObject/009446.png
VOCdevkit/VOC2007/SegmentationObject/009458.png
VOCdevkit/VOC2007/SegmentationObject/009464.png

VOCdevkit/VOC2007/SegmentationObject/009527.png
VOCdevkit/VOC2007/SegmentationObject/009533.png
VOCdevkit/VOC2007/SegmentationObject/009550.png
VOCdevkit/VOC2007/SegmentationObject/009562.png
VOCdevkit/VOC2007/SegmentationObject/009580.png
VOCdevkit/VOC2007/SegmentationObject/009597.png
VOCdevkit/VOC2007/SegmentationObject/009605.png
VOCdevkit/VOC2007/SegmentationObject/009618.png
VOCdevkit/VOC2007/SegmentationObject/009649.png
VOCdevkit/VOC2007/SegmentationObject/009654.png
VOCdevkit/VOC2007/SegmentationObject/009655.png
VOCdevkit/VOC2007/SegmentationObject/009684.png
VOCdevkit/VOC2007/SegmentationObject/009687.png
VOCdevkit/VOC2007/SegmentationObject/009691.png
VOCdevkit/VOC2007/SegmentationObject/009706.png
VOCdevkit/VOC2007/SegmentationObject/009709.png
VOCdevkit/VOC2007/SegmentationObject/009724.png
VOCdevkit/VOC2007/SegmentationObject/009756.png
VOCdevkit/VOC2007/SegmentationObject/009764.png
VOCdevkit/VOC2007/SegmentationObject/009794.png
VOCdevkit/VOC2007/SegmentationObject/009807.png
VOCdevkit/VOC2007/SegmentationObject/009832.png
VOCdevkit/VOC2007/SegmentationObject/009841.png
VOCdevkit/VOC2007/SegmentationObject/009897.png
VOCdevkit/VOC2007/SegmentationObject/009911.png
VOCdevkit/VOC2007/SegmentationObject/009923.png
VOCdevkit/VOC2007/SegmentationObject/009938.png
VOCdevkit/VOC2007/SegmentationObject/009947.png
VOCdevkit/VOC2007/SegmentationObject/009950.png

```

```
!mv VOCdevkit datasets
```

```
from detectron2.engine import DefaultTrainer
```

```

cfg = get_cfg()
cfg.merge_from_file(model_zoo.get_config_file("COCO-Detection/retinanet_R_50_FPN_3x.yaml"))
cfg.OUTPUT_DIR = 'MyVOCTraining'
cfg.DATASETS.TRAIN = ("voc_2007_train",)
cfg.DATASETS.TEST = ( )

```

```

cfg.DATALOADER.NUM_WORKERS = 1
cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url("COCO-Detection/retinanet_R_50_FPN_3x.yaml")
cfg.SOLVER.IMS_PER_BATCH = 1
cfg.SOLVER.BASE_LR = 0.0001 # pick a good LR
cfg.SOLVER.MAX_ITER = 3000
cfg.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE = 128
cfg.MODEL.ROI_HEADS.NUM_CLASSES = 20

os.makedirs(cfg.OUTPUT_DIR, exist_ok=True)
trainer = DefaultTrainer(cfg)
trainer.resume_or_load(resume=False)
trainer.train()

```

```

[11/18 20:33:13 d2.utils.events]: eta: 0:14:08 iter: 1919 total_loss: 0.5091 loss
[11/18 20:33:28 d2.utils.events]: eta: 0:13:52 iter: 1939 total_loss: 0.2309 loss
[11/18 20:33:44 d2.utils.events]: eta: 0:13:35 iter: 1959 total_loss: 0.3787 loss
[11/18 20:33:59 d2.utils.events]: eta: 0:13:19 iter: 1979 total_loss: 0.4733 loss
[11/18 20:34:15 d2.utils.events]: eta: 0:13:04 iter: 1999 total_loss: 0.2636 loss
[11/18 20:34:30 d2.utils.events]: eta: 0:12:48 iter: 2019 total_loss: 0.3307 loss
[11/18 20:34:46 d2.utils.events]: eta: 0:12:31 iter: 2039 total_loss: 0.3108 loss
[11/18 20:35:01 d2.utils.events]: eta: 0:12:15 iter: 2059 total_loss: 0.3146 loss
[11/18 20:35:16 d2.utils.events]: eta: 0:12:00 iter: 2079 total_loss: 0.4217 loss
[11/18 20:35:32 d2.utils.events]: eta: 0:11:45 iter: 2099 total_loss: 0.2993 loss
[11/18 20:35:47 d2.utils.events]: eta: 0:11:28 iter: 2119 total_loss: 0.3754 loss
[11/18 20:36:03 d2.utils.events]: eta: 0:11:11 iter: 2139 total_loss: 0.3669 loss
[11/18 20:36:19 d2.utils.events]: eta: 0:10:54 iter: 2159 total_loss: 0.5506 loss
[11/18 20:36:34 d2.utils.events]: eta: 0:10:41 iter: 2179 total_loss: 0.416 loss
[11/18 20:36:50 d2.utils.events]: eta: 0:10:26 iter: 2199 total_loss: 0.2713 loss
[11/18 20:37:05 d2.utils.events]: eta: 0:10:10 iter: 2219 total_loss: 0.2383 loss
[11/18 20:37:20 d2.utils.events]: eta: 0:09:53 iter: 2239 total_loss: 0.2838 loss
[11/18 20:37:36 d2.utils.events]: eta: 0:09:37 iter: 2259 total_loss: 0.4469 loss
[11/18 20:37:52 d2.utils.events]: eta: 0:09:22 iter: 2279 total_loss: 0.2978 loss
[11/18 20:38:08 d2.utils.events]: eta: 0:09:06 iter: 2299 total_loss: 0.3353 loss
[11/18 20:38:25 d2.utils.events]: eta: 0:08:52 iter: 2319 total_loss: 0.4355 loss
[11/18 20:38:41 d2.utils.events]: eta: 0:08:37 iter: 2339 total_loss: 0.2347 loss
[11/18 20:38:58 d2.utils.events]: eta: 0:08:22 iter: 2359 total_loss: 0.3628 loss
[11/18 20:39:12 d2.utils.events]: eta: 0:08:06 iter: 2379 total_loss: 0.4204 loss
[11/18 20:39:28 d2.utils.events]: eta: 0:07:51 iter: 2399 total_loss: 0.2305 loss
[11/18 20:39:44 d2.utils.events]: eta: 0:07:36 iter: 2419 total_loss: 0.2895 loss
[11/18 20:39:59 d2.utils.events]: eta: 0:07:19 iter: 2439 total_loss: 0.3893 loss
[11/18 20:40:14 d2.utils.events]: eta: 0:07:04 iter: 2459 total_loss: 0.3144 loss
[11/18 20:40:29 d2.utils.events]: eta: 0:06:48 iter: 2479 total_loss: 0.354 loss
[11/18 20:40:45 d2.utils.events]: eta: 0:06:33 iter: 2499 total_loss: 0.3286 loss
[11/18 20:41:02 d2.utils.events]: eta: 0:06:18 iter: 2519 total_loss: 0.384 loss
[11/18 20:41:18 d2.utils.events]: eta: 0:06:02 iter: 2539 total_loss: 0.2953 loss
[11/18 20:41:33 d2.utils.events]: eta: 0:05:47 iter: 2559 total_loss: 0.4076 loss
[11/18 20:41:49 d2.utils.events]: eta: 0:05:31 iter: 2579 total_loss: 0.3293 loss
[11/18 20:42:04 d2.utils.events]: eta: 0:05:15 iter: 2599 total_loss: 0.2555 loss

[11/18 20:42:20 d2.utils.events]: eta: 0:05:00 iter: 2619 total_loss: 0.2009 loss
[11/18 20:42:35 d2.utils.events]: eta: 0:04:43 iter: 2639 total_loss: 0.2147 loss
[11/18 20:42:51 d2.utils.events]: eta: 0:04:28 iter: 2659 total_loss: 0.2838 loss
[11/18 20:43:06 d2.utils.events]: eta: 0:04:12 iter: 2679 total_loss: 0.3142 loss
[11/18 20:43:21 d2.utils.events]: eta: 0:03:56 iter: 2699 total_loss: 0.2127 loss
[11/18 20:43:37 d2.utils.events]: eta: 0:03:40 iter: 2719 total_loss: 0.304 loss
[11/18 20:43:53 d2.utils.events]: eta: 0:03:24 iter: 2739 total_loss: 0.2223 loss
[11/18 20:44:08 d2.utils.events]: eta: 0:03:09 iter: 2759 total_loss: 0.315 loss

```

```

[11/18 20:44:08 d2.utils.events]: eta: 0:03:09 iter: 2759 total_loss: 0.313 loss
[11/18 20:44:24 d2.utils.events]: eta: 0:02:53 iter: 2779 total_loss: 0.2391 loss
[11/18 20:44:39 d2.utils.events]: eta: 0:02:37 iter: 2799 total_loss: 0.4163 loss
[11/18 20:44:55 d2.utils.events]: eta: 0:02:21 iter: 2819 total_loss: 0.2607 loss
[11/18 20:45:11 d2.utils.events]: eta: 0:02:06 iter: 2839 total_loss: 0.2613 loss
[11/18 20:45:27 d2.utils.events]: eta: 0:01:50 iter: 2859 total_loss: 0.2303 loss
[11/18 20:45:42 d2.utils.events]: eta: 0:01:34 iter: 2879 total_loss: 0.294 loss
[11/18 20:45:58 d2.utils.events]: eta: 0:01:18 iter: 2899 total_loss: 0.1452 loss
[11/18 20:46:13 d2.utils.events]: eta: 0:01:03 iter: 2919 total_loss: 0.2154 loss
[11/18 20:46:28 d2.utils.events]: eta: 0:00:47 iter: 2939 total_loss: 0.3519 loss
[11/18 20:46:44 d2.utils.events]: eta: 0:00:31 iter: 2959 total_loss: 0.2633 loss
[11/18 20:47:00 d2.utils.events]: eta: 0:00:15 iter: 2979 total_loss: 0.3403 loss
[11/18 20:47:16 d2.utils.events]: eta: 0:00:00 iter: 2999 total_loss: 0.249 loss
[11/18 20:47:17 d2.engine.hooks]: Overall training speed: 2998 iterations in 0:39:05
[11/18 20:47:17 d2.engine.hooks]: Total training time: 0:39:07 (0:00:02 on hooks)

```

```

# Look at training curves in tensorboard:
%reload_ext tensorboard
%tensorboard --logdir MyVOCTraining/

```

TensorBoard

SCALARS

TIME SERIES

INACTIVE

☐ Show data download links

Filter tags (regular expressions supported)

☐ Ignore outliers in chart scalingTooltip sorting
method:

default ▼

data_time ▼

eta_seconds ▼

```
# Inference should use the config with parameters that are used in training
# cfg now already contains everything we've set previously. We changed it a little bit for inference
cfg.MODEL.WEIGHTS = os.path.join(cfg.OUTPUT_DIR, "model_final.pth") # path to the model weights
cfg.MODEL.ROI_HEADS.SCORE_THRESH_TEST = 0.7 # set a custom testing threshold
predictor = DefaultPredictor(cfg)
```

```
from detectron2.evaluation import PascalVOCDetectionEvaluator, inference_on_dataset
from detectron2.data import build_detection_test_loader
evaluator = PascalVOCDetectionEvaluator("voc_2007_val")
val_loader = build_detection_test_loader(cfg, "voc_2007_val")
print(inference_on_dataset(predictor.model, val_loader, evaluator))
```

```
[11/18 20:49:52 d2.data.build]: Distribution of instances among all 20 categories:
```

category	#instances	category	#instances	category	#instances
aeroplane	175	bicycle	216	bird	305
boat	190	bottle	296	bus	141
car	818	cat	198	chair	706
cow	171	diningtable	162	dog	267
horse	199	motorbike	197	person	2742
pottedplant	320	sheep	162	sofa	207
train	170	tvmonitor	176		
total	7818				

```
[11/18 20:49:52 d2.data.dataset_mapper]: [DatasetMapper] Augmentations used in inference
```

```
[11/18 20:49:52 d2.data.common]: Serializing 2510 elements to byte tensors and concatenating them all
```

```
[11/18 20:49:52 d2.data.common]: Serialized dataset takes 1.14 MiB
```

```
[11/18 20:49:52 d2.evaluation.evaluator]: Start inference on 2510 batches
```

```
/usr/local/lib/python3.7/dist-packages/detectron2/structures/image_list.py:88: UserWarning: max_size = (max_size + (stride - 1)) // stride * stride
```

```
[11/18 20:49:56 d2.evaluation.evaluator]: Inference done 11/2510. Dataloading: 0.001
```

```
[11/18 20:50:02 d2.evaluation.evaluator]: Inference done 25/2510. Dataloading: 0.002
```

```
[11/18 20:50:07 d2.evaluation.evaluator]: Inference done 39/2510. Dataloading: 0.002
```

```
[11/18 20:50:12 d2.evaluation.evaluator]: Inference done 53/2510. Dataloading: 0.002
```

```
[11/18 20:50:17 d2.evaluation.evaluator]: Inference done 67/2510. Dataloading: 0.002
```

```
[11/18 20:50:22 d2.evaluation.evaluator]: Inference done 81/2510. Dataloading: 0.002
```

```
[11/18 20:50:27 d2.evaluation.evaluator]: Inference done 94/2510. Dataloading: 0.002
```

```
[11/18 20:50:32 d2.evaluation.evaluator]: Inference done 108/2510. Dataloading: 0.002
```

```
[11/18 20:50:37 d2.evaluation.evaluator]: Inference done 122/2510. Dataloading: 0.002
```

```
[11/18 20:50:42 d2.evaluation.evaluator]: Inference done 135/2510. Dataloading: 0.002
```

```
[11/18 20:50:48 d2.evaluation.evaluator]: Inference done 148/2510. Dataloading: 0.002
```



```

[11/18 20:50:53 d2.evaluation.evaluator]: Inference done 162/2510. Dataloading: 0.00
[11/18 20:50:58 d2.evaluation.evaluator]: Inference done 176/2510. Dataloading: 0.00
[11/18 20:51:03 d2.evaluation.evaluator]: Inference done 190/2510. Dataloading: 0.00
[11/18 20:51:08 d2.evaluation.evaluator]: Inference done 203/2510. Dataloading: 0.00
[11/18 20:51:14 d2.evaluation.evaluator]: Inference done 217/2510. Dataloading: 0.00
[11/18 20:51:19 d2.evaluation.evaluator]: Inference done 231/2510. Dataloading: 0.00
[11/18 20:51:24 d2.evaluation.evaluator]: Inference done 245/2510. Dataloading: 0.00
[11/18 20:51:29 d2.evaluation.evaluator]: Inference done 259/2510. Dataloading: 0.00
[11/18 20:51:35 d2.evaluation.evaluator]: Inference done 273/2510. Dataloading: 0.00
[11/18 20:51:40 d2.evaluation.evaluator]: Inference done 286/2510. Dataloading: 0.00
[11/18 20:51:45 d2.evaluation.evaluator]: Inference done 300/2510. Dataloading: 0.00
[11/18 20:51:50 d2.evaluation.evaluator]: Inference done 314/2510. Dataloading: 0.00
[11/18 20:51:55 d2.evaluation.evaluator]: Inference done 327/2510. Dataloading: 0.00
[11/18 20:52:00 d2.evaluation.evaluator]: Inference done 340/2510. Dataloading: 0.00
[11/18 20:52:05 d2.evaluation.evaluator]: Inference done 354/2510. Dataloading: 0.00
[11/18 20:52:10 d2.evaluation.evaluator]: Inference done 367/2510. Dataloading: 0.00
[11/18 20:52:15 d2.evaluation.evaluator]: Inference done 381/2510. Dataloading: 0.00
[11/18 20:52:21 d2.evaluation.evaluator]: Inference done 395/2510. Dataloading: 0.00
[11/18 20:52:26 d2.evaluation.evaluator]: Inference done 409/2510. Dataloading: 0.00
[11/18 20:52:31 d2.evaluation.evaluator]: Inference done 423/2510. Dataloading: 0.00
[11/18 20:52:36 d2.evaluation.evaluator]: Inference done 437/2510. Dataloading: 0.00
[11/18 20:52:42 d2.evaluation.evaluator]: Inference done 451/2510. Dataloading: 0.00
[11/18 20:52:47 d2.evaluation.evaluator]: Inference done 465/2510. Dataloading: 0.00
[11/18 20:52:52 d2.evaluation.evaluator]: Inference done 479/2510. Dataloading: 0.00
[11/18 20:52:57 d2.evaluation.evaluator]: Inference done 493/2510. Dataloading: 0.00
[11/18 20:53:03 d2.evaluation.evaluator]: Inference done 507/2510. Dataloading: 0.00
[11/18 20:53:08 d2.evaluation.evaluator]: Inference done 521/2510. Dataloading: 0.00
[11/18 20:53:13 d2.evaluation.evaluator]: Inference done 534/2510. Dataloading: 0.00
[11/18 20:53:18 d2.evaluation.evaluator]: Inference done 548/2510. Dataloading: 0.00

```

```

dataset_dicts = detectron2.data.get_detection_dataset_dicts('voc_2007_val')
for d in random.sample(dataset_dicts, 3):
    im = cv2.imread(d["file_name"])
    outputs = predictor(im) # format is documented at https://detectron2.readthedocs.io/tutorials/visualization.html
    v = Visualizer(im[:, :, ::-1],
                   detectron2.data.MetadataCatalog.get('voc_2007_val'),
                   scale=0.5,
                   )
    out = v.draw_instance_predictions(outputs["instances"].to("cpu"))
    cv2.imshow(out.get_image()[:, :, ::-1])

```



```
[11/18 21:17:33 d2.data.build]: Removed 0 images with no usable annotations. 2510 image
/usr/local/lib/python3.7/dist-packages/detectron2/structures/image_list.py:88: UserWarn
    max_size = (max_size + (stride - 1)) // stride * stride
```



✓ 3s completed at 1:17 PM

